arXiv:2610.19934v1 [cs.CL] 14 Oct 2026

To Reason but Not to Stop: Understanding and Recovering Pruned Large Reasoning Models

Minh-Anh Nguyen[1] **Nguyen T. Ha Lan**[2] **Xiyan Fu**[1] **Dung D. Le**[2] **Wei Lu**[1]
[1]Nanyang Technological University, Singapore
[2]VinUniversity, Vietnam

Abstract

Weight pruning can degrade the accuracy of large reasoning models (LRMs), yet it remains unclear whether it damages reasoning itself. We find that a substantial fraction of this degradation stems from a failure to *stop* rather than to *reason*. As sparsity increases, pruned LRMs increasingly fail to terminate within the token budget, often after the correct answer has appeared in their reasoning, instead falling into repetitive loops, even though they still recognize where the dense model would stop. We further find that the hidden-state difference between terminating and looping continuations is consistent across pruned checkpoints, suggesting a shared direction associated with committing to an answer. Building on this, we propose Termination-Guided Activation Steering (TGAS), a training-free method that adds this direction to the residual stream at inference time. Across six LRMs, three pruning pipelines, and math and code benchmarks, TGAS improves termination and overall accuracy while reducing repetitive reasoning and generation length, showing that this failure is actionable.

1 Introduction

Large reasoning models (LRMs) achieve strong performance by generating long chain-of-thought traces, but this makes inference expensive (Zhang et al., 2026a; Guo et al., 2025). Weight pruning offers a natural approach to reduce cost by zeroing out selected parameters (Frantar and Alistarh, 2023; Sun et al., 2024). However, as consistently observed across pruning methods, pruning LRMs can substantially degrade accuracy on reasoning tasks (Ji et al., 2025; Lucas et al., 2026; Zhang et al., 2026b). Existing efforts have primarily focused on mitigating this degradation through improved pruning calibration (Ji et al., 2025; Lucas et al., 2026; Zhang et al., 2026b), alongside post-pruning recovery via retraining or fine-tuning (Liu et al., 2025b; Muralidharan et al., 2024; Kim et al., 2024; Xia et al., 2024a). Yet, it remains unclear whether this degradation reflects a loss of reasoning capability or a different kind of failure.

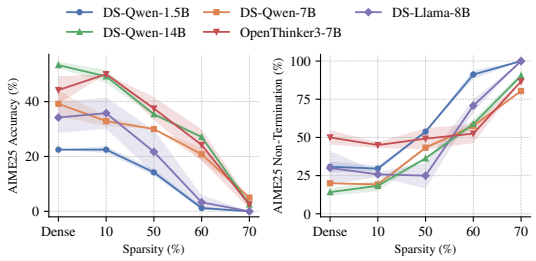

Figure 1: Accuracy (left) and non-termination rate (right) on AIME-25 across pruning settings (SparseGPT with OBC calibration, 16,384-token budget, averaged over 4 runs). Accuracy decreases as non-termination becomes more frequent.

We find that much of this degradation stems from a failure to *stop* rather than a failure to *reason*. As sparsity increases, pruned LRMs increasingly fail to terminate and commit to a final answer within the token budget, and accuracy falls accordingly (Figure 1). Strikingly, many of these non-terminating traces already contain the correct answer, yet the model keeps reasoning until the budget is exhausted. We refer to this behavior as ***reasoning divergence***. On problems solved by the dense model, pruned models remain largely accurate whenever they terminate, and non-termination becomes the dominant source of errors as sparsity grows (Section 3.1).

Why do pruned models fail to stop? Pruning does not erase the stopping signal: when replaying dense traces, pruned models still assign high probability to the end-of-thinking token where the dense model stops. Instead, during free-running generation they increasingly revisit earlier reasoning rather than moving forward, and fall into repetitive loops. Moreover, the hidden-state difference between continuations that terminate and

those that loop points in a similar direction across pruned checkpoints, suggesting that the shift between looping and committing to an answer is captured by a shared direction in representation space.

Based on this diagnosis, we propose *Termination-Guided Activation Steering* (TGAS), which adds this direction to the residual stream at inference time to steer pruned models toward termination without retraining or weight updates. TGAS consistently improves termination and accuracy across six LRMs, three pruning pipelines, and math and code benchmarks by reducing repetitive reasoning. Our contributions are threefold: (i) We show that a substantial part of pruning-induced degradation in LRMs arises from reasoning divergence, a failure to commit rather than a loss of reasoning ability; (ii) We identify a hidden-state direction that separates terminating from looping continuations, and build on it TGAS, a training-free steering method; (iii) We show that TGAS consistently improves termination and overall accuracy across models, pruning methods, and benchmarks by reducing looping.

2 Background and Related Work

Pruning for LRMs. One-shot unstructured pruning formulates compression as layer-wise reconstruction (Frantar et al., 2022): given a weight matrix $\mathbf{W}_\ell$ and calibration inputs $\mathbf{X}_\ell$, it solves $\min_{\|\hat{\mathbf{W}}_\ell\|_0 \le S} \|(\mathbf{W}_\ell - \hat{\mathbf{W}}_\ell)\mathbf{X}_\ell\|_F^2$, where S is the number of retained weights. We focus on two retraining-free methods: SparseGPT (Frantar and Alistarh, 2023), which uses second-order information to compensate for pruning errors, and Wanda (Sun et al., 2024), which scores weights by magnitude and input-activation norm. To better capture long reasoning traces, recent LRM pruning methods improve calibration via on-policy chain-of-thought RAC (Lucas et al., 2026), reasoning-step importance OBC (Nguyen et al., 2026), or self-generated traces (Ji et al., 2025; Wang et al., 2025; Xiang et al., 2025), detailed preliminaries in Appendix B.

Related Works. LRMs can overthink and enter repetitive loops (Chen et al., 2026; Pipis et al., 2026), which quantization can further prolong (Lotfi et al., 2026). We show that pruning amplifies looping while largely preserving the stopping signal, distinguishing failures to stop from failures to reason. Activation steering offers a way to modify behavior by adding directions to the residual stream, with difference-in-means directions shown to control behaviors of LLMs (Arditi et al., 2024). For reasoning, SEAL (Chen et al., 2025) suppresses reflection, while Plan-CoT (Sheng et al., 2025) modulates reasoning length. In contrast, TGAS targets pruning-induced non-termination using contrasting continuations of the same prefix, yielding a direction weakly aligned with SEAL's (Appendix F.8).

3 Understanding Pruning-Induced Reasoning Degradation

We investigate three questions: (i) Do pruned models lose reasoning ability or only the ability to stop? (ii) Is non-termination caused by limited generation budgets? (iii) Why do models fail to stop despite reaching the correct answer? We address the first two questions in Section 3.1 and the third in Section 3.2. Recall that we refer to this phenomenon as *reasoning divergence*, characterized by *non-termination*, driven by repetitive *looping* and an *failure to commit* to final answers.

Setting. We evaluate DeepSeek-R1-Distill-Qwen (DS-Qwen), DeepSeek-R1-Distill-Llama (DS-Llama) (Guo et al., 2025), and OpenThinker3 (OT3) (Guha et al., 2026) under one-shot unstructured pruning of decoder linear layers at varying sparsity levels, using SparseGPT or Wanda with OBC, RAC, or C4 calibration. Benchmarks include MATH-500 Levels 4-5 (hereafter MATH-500) (Lightman et al., 2024), AIME-25 (Zhang and Math-AI, 2025), and LiveCodeBench-v6 (LCB) (Jain et al., 2025); OpenR1-Math-220K (OpenR1-Math) (Hugging Face, 2025) is used only for calibration and steering-direction construction. We use a 16,384-token budget, $T = 0.6$, and top-$p = 0.95$, averaging $k = 4$ responses on three benchmarks. Our default uses SparseGPT with OBC calibration on model-generated reasoning traces, counting a response as correct only if it commits to a final answer, detailed in Appendix C.

3.1 Do Pruned Models Lose the Ability to Reason?

Accuracy on Dense-Correct Problems. To determine whether reasoning divergence reflects a loss of reasoning capability, we evaluate pruned models on problems correctly solved by their

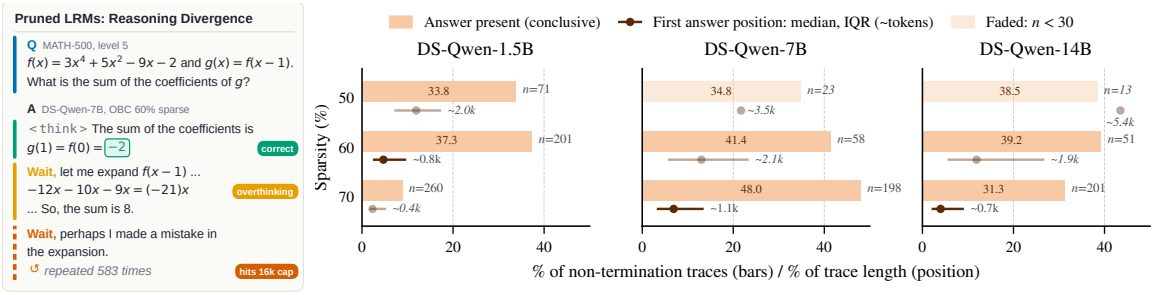

Figure 2: Gold answers in non-terminating reasoning traces of DeepSeek models on MATH-500. Left: an example of reasoning divergence, where the pruned model continues reasoning after reaching the correct answer. Right: the fraction of traces containing the gold answer and the median position of its first occurrence. n denotes the number of non-terminating samples.

Model	Ckpt	Kept	No-Term.	Wrong	Kept|term
DS-Qwen -7B	obc_50	91.7±0.8	3.8±0.3	4.5±0.6	95.3±0.6
	obc_60	77.8±0.7	16.8±0.1	5.4±0.6	93.5±0.7
DS-Qwen -14B	obc_50	89.2±0.6	2.4±1.5	8.4±0.9	91.4±0.8
	obc_60	74.3±2.8	16.4±1.9	9.3±0.9	88.9±1.4
OT3-7B	obc_50	79.3±2.9	13.1±2.4	7.6±1.9	91.3±2.2
	obc_60	63.6±3.4	29.8±3.3	6.6±1.8	90.6±2.5

Table 1: Accuracy on MATH-500 problems correctly solved by the dense model. **Kept**: correct; **No-Term.**: reasoning fails to terminate; **Wrong**: terminated but incorrect; **Kept|term**: accuracy among terminated generations. From 50% to 60% sparsity, non-termination grows sharply while the terminated-but-wrong rate stays nearly flat, so accuracy among terminated generations remains high.

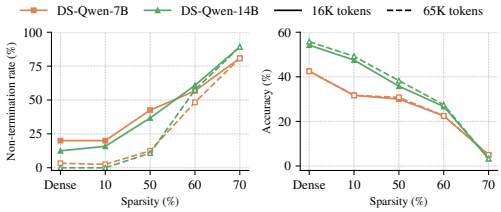

Figure 3: Non-termination rate and accuracy of DeepSeek models on AIME-25 under 16K- and 65K-token budgets.

dense counterparts. Table 1 shows that as sparsity grows, the additional errors come mainly from non-termination rather than from incorrect answers. On MATH-500, going from 50% to 60% sparsity raises the non-termination rate from 3.8% to 16.8% for DS-Qwen-7B, and from 2.4% to 16.4% for the 14B counterpart, while the terminated-but-wrong rate stays nearly unchanged. Terminated generations largely preserve their accuracy at high sparsity, despite fewer dense-correct samples remaining correct. These results suggest that substantial reasoning capability remains intact under moderate pruning, while termination becomes the primary bottleneck as sparsity grows (full results in Appendix D.1).

Larger Generation Budget. We ask whether reasoning divergence is simply caused by an insufficient generation budget. As shown in Figure 3, increasing the budget from 16,384 to 65,536 tokens substantially reduces non-termination at low and moderate sparsity, e.g., from 42.5% to 12.5% for 7B at 50% sparsity, but leaves accuracy nearly unchanged for DS-Qwen 7B and 14B, while at 60–70% sparsity non-termination barely decreases. Appendix D.2 shows that the larger budget recovers only 0–3 additional correct traces across settings, suggesting that pruning-induced failures are not primarily due to insufficient generation length.

Correct Answer Presence in Reasoning Traces. Since a larger budget rarely helps, we next ask what the non-terminating traces contain: does the reasoning already reach the correct answer even when the model fails to terminate and finalize it? To answer this question, we use text-based rules, termed *Conclusive*, that detect answer-like statements following cues such as *"thus the answer"* or *"so we get"*(Appendix C.2).[1] Figure 2 shows that a substantial fraction of non-terminating traces already contain the correct answer: at 60% sparsity, this holds for 41.4% and 39.2% of non-terminating traces for DS-Qwen 7B and 14B, respectively, with the answer appearing early at a median of 13.0% and 11.9% of the trace length, roughly 2.1k and 1.9k tokens. Despite reaching the correct answer, models continue generating to the token limit, often entering repetitive or corrective reasoning. This suggests that non-termination often reflects reasoning divergence rather than an inability to reach the correct answer. We provide a headroom analysis in Appendix E.3, quantifying the recoverable accuracy loss by counting non-terminating traces that already contain the correct

[1] We validate *Conclusive* against an independent LLM judge: on 250 sampled traces, the two agree on 96.0% of cases ($\kappa = 0.898$; Appendix C.3).

		DS-Qwen-7B			DS-Qwen-14B		
		obc_50	obc_60	obc_70	obc_50	obc_60	obc_70
AIME-25	$p(\texttt{</think>})$	0.998	0.957	0.036	0.995	0.854	0.073
	Med. rank	0	0	5.5	0	0	2
MATH-500	$p(\texttt{</think>})$	0.996	0.924	0.040	0.998	0.907	0.130
	Med. rank	0	0	7	0	0	3

Table 2: Teacher-forced stopping signal at the dense model's STOP position: probability of </think> and its median rank (0 means </think> is the top prediction).

	DS-Qwen-7B				OT3-7B			
	dense	obc_50	obc_60	obc_70	dense	obc_50	obc_60	obc_70
n	794	788	796	782	800	800	796	790
Δ_{cyc}	−15.3	−11.9	−9.6	−5.0	−23.1	−15.8	−13.2	−5.6
Entropy	1.55	1.51	1.41	1.44	1.03	1.22	1.33	1.60

Table 3: Log-probability ratio between cyclic and progressive continuations at decision points on AIME-25. n: number of decision points; **Entropy**: predictive entropy at the decision points. Less negative (larger) Δ_{cyc} values indicate a stronger preference for the repetitive continuation.

answer as correct.

3.2 Why Do Pruned Models Fail to Terminate?

Section 3.1 shows that pruned LRMs largely retain the ability to reach correct answers but increasingly fail to terminate. We show that the stopping signal remains intact, yet during free-running generation, models struggle to commit to final answers and fall into repetitive reasoning loops.

Does pruning damage the stopping mechanism? We test whether pruning impairs stopping recognition via teacher forcing. Let $\mathcal{D}_{\text{cov}}$ denote dense traces that terminate correctly. We replay these traces through each pruned checkpoint and measure $p(\texttt{</think>})$ and its rank at the dense model's termination position (**STOP**). As shown in Table 2, at 50–60% sparsity, the signal remains strong: for 7B, $p(\texttt{</think>}) > 0.92$ with median rank 0 on both AIME-25 and MATH-500, while 14B exceeds 0.85 with the same rank. At 70%, it drops to 0.036–0.040 for 7B and 0.073–0.130 for 14B. Comparing STOP with positions more than 50 tokens earlier, we find that </think> remains substantially more likely at STOP at all sparsities, but at 70% its log-probability falls below that of the ordinarily generated tokens, detailed in Appendix D.3. Thus, the stopping signal remains largely intact under moderate pruning, suggesting that the main failure is reaching, rather than recognizing, the stopping point. At 70% sparsity, the signal remains localized but weakens, consistent with severe pruning also impairing stopping recognition.

Definition. We investigate whether pruning alters reasoning trajectories at moments of reconsideration, where a model may either move forward or revisit previous reasoning. We define a *decision point* as the start of a paragraph beginning with a predefined reconsideration cue (e.g., *"wait"*, *"hmm"*, full list in Appendix D.4), following Lotfi et al. (2026). We label each paragraph CYCLIC if its 8-word n-gram set has Jaccard similarity ≥ 0.5 with any earlier paragraph in the same trace, and PROGRESS otherwise, following Pipis et al. (2026), as illustrated in Figure 4.

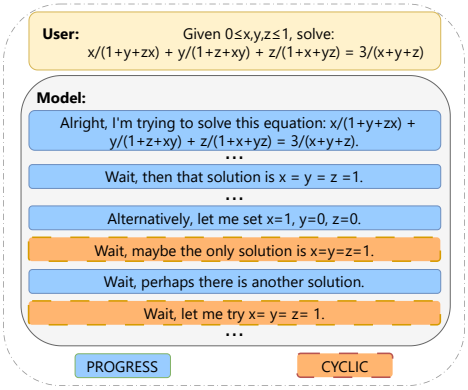

Figure 4: PROGRESS and CYCLIC behaviors illustration.

Teacher-Forced Preferences. To measure whether a model prefers to repeat or to move forward, we compare two candidate next paragraphs at each decision point of a dense trace: repeating the preceding paragraph (CYCLIC) or continuing with the next paragraph that adds new content (PROGRESS). We score the first 32 tokens of each candidate under teacher forcing with the same prefix, and reuse the same decision points and candidates across all pruning checkpoints. We restrict evaluation to decision points that are followed by a PROGRESS paragraph, and define:

$$\Delta_{\text{cyc}} = \sum_{t=1}^{32} \Big[\log p_\theta\big(\mathbf{x}_t^{\text{cyclic}} \mid \text{prefix}, \mathbf{x}_{<t}^{\text{cyclic}}\big) - \log p_\theta\big(\mathbf{x}_t^{\text{progress}} \mid \text{prefix}, \mathbf{x}_{<t}^{\text{progress}}\big) \Big]$$

where a larger Δ_{cyc} indicates a stronger relative preference for CYCLIC over PROGRESS.

As shown in Table 3, pruning progressively shifts the model's preference from progressive toward cyclic continuations at decision points; Appendix D.4 shows a similar shift at randomly sam-

	DS-Qwen-7B				DS-Qwen-14B			
	dense	obc_50	obc_60	obc_70	dense	obc_50	obc_60	obc_70
P(cyc | cyc)	0.029	0.283	0.327	0.307	0.000	0.203	0.271	0.324
P(cyc | prog)	0.026	0.043	0.064	0.085	0.029	0.033	0.059	0.081
Count (k)	1.3	51.5	85.7	113.4	0.8	47.3	68.3	136.8

Table 4: Transition dynamics between CYCLIC and PROGRESS paragraphs on AIME-25 from shared branch points. $P(\text{cyc} \mid \text{cyc})$ measures loop persistence, while $P(\text{cyc} \mid \text{prog})$ measures loop entry; **Count**: number of adjacent paragraph pairs, in thousands; dense rows have fewer pairs due to early termination, while pruned checkpoints retain reasoning divergence despite using dense correct traces.

pled paragraph, indicating a general preference for repetition rather than one specific to reconsideration cues. For DS-Qwen-7B, Δ_{cyc} increases from -15.3 for the dense model to -5.0 at 70% sparsity; OT3-7B shows a similar shift from -23.1 to -5.6, other models in Appendix D.4. Predictive entropy changes little for DS-Qwen-7B but increases for OT3-7B, whereas Δ_{cyc} shifts consistently for both, suggesting that the shift is not merely a by-product of increased uncertainty. Overall, pruning makes models increasingly favor repetitive over progressive reasoning, which we confirm in free-running generation below.

Free-Running Continuations. The preceding analysis shows that pruning shifts model behavior toward repetitive loops rather than progressing new reasoning steps. We next examine how pruned models continue dense reasoning traces. Following Section 3.1, we revisit the *Conclusive* point. For each trace, we define the branch point b as the first eligible paragraph start after *Conclusive* such that the prefix $\mathbf{x}_{<b}$ is sufficient to identify the correct answer. We teacher-force the prefix through the pruned checkpoint and sample N continuations, discarding traces without eligible branch points. Thus, continuations share the same prefix, checkpoint, and problem, differing only beyond the branch point.

Table 4 reveals a shift toward cyclic reasoning in pruned checkpoints. Both loop persistence, $P(\text{cyc} \mid \text{cyc})$, and loop entry, $P(\text{cyc} \mid \text{prog})$, generally increase with pruning. For DS-Qwen 7B and 14B at 70% sparsity, $P(\text{cyc} \mid \text{cyc})$ reaches 0.307 and 0.324, while $P(\text{cyc} \mid \text{prog})$ reaches 0.085 and 0.081, respectively, compared with at most 0.029 for the dense models (other models in Appendix D.4). This indicates that pruning is associated with both more persistent cyclic reasoning and more frequent transitions from progressive to cyclic states.

4 Methodology

If looping reflects a failure to commit rather than a loss of capability, a pruned model should be steerable toward committing without weight updates. We use a hidden-state direction that separates terminating from looping continuations under the same prefix.

4.1 Constructing the Termination Direction

We reuse the dense correct traces $\mathcal{D}_{\text{cov}}$ and branch points b from Section 3.2; at each branch point, the prefix $\mathbf{x}_{<b}$ already suffices to derive the correct answer. From each prefix, we sample N continuations with the pruned checkpoint and split them into $\mathcal{C}_b^+$, continuations that emit `</think>` within the token budget, and $\mathcal{C}_b^-$, continuations that reach the budget without terminating (N, set sizes and median lengths are reported in Appendix C.4.5). For each layer $\ell \in \{1, \ldots, L\}$, we represent a continuation c by its mean hidden state over the span from the branch point b to the dense trace's `</think>` position t_{end}:

$$\mathbf{s}_\ell(c) = \frac{1}{t_{\text{end}} - b} \sum_{t=b}^{t_{\text{end}}-1} \mathbf{h}_{\ell,t}(c),$$

For continuations that end before t_{end}, we average hidden states from b to the continuation's end. We use the dense trace's t_{end} as a common reference to exclude the looping tail from the mean and ensure comparable representations. We compute the class means as:

$$\boldsymbol{\mu}_{\ell,b}^+ = \frac{1}{|\mathcal{C}_b^+|} \sum_{c \in \mathcal{C}_b^+} \mathbf{s}_\ell(c), \quad \boldsymbol{\mu}_{\ell,b}^- = \frac{1}{|\mathcal{C}_b^-|} \sum_{c \in \mathcal{C}_b^-} \mathbf{s}_\ell(c).$$

Following the difference-in-means approach of Arditi et al. (2024), we then average the class difference over branch points:

$$\mathbf{r}_\ell = \frac{1}{|\mathcal{B}|} \sum_{b \in \mathcal{B}} (\boldsymbol{\mu}_{\ell,b}^+ - \boldsymbol{\mu}_{\ell,b}^-), \tag{1}$$

where $\mathcal{B}$ denotes the set of eligible branch points collected from $\mathcal{D}_{\text{cov}}$. Because the difference is taken within each branch point, the two classes share the same problem and prefix, so problem-specific content cancels and $\mathbf{r}_\ell$ isolates the contrast between committing and looping. This yields a candidate direction $\mathbf{r}_\ell$ for each layer. Since layers lie in different residual spaces, we compare two

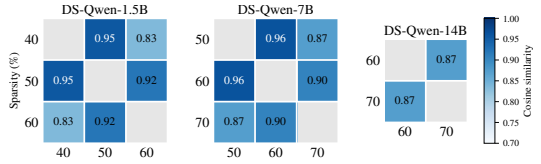

Figure 5: Cosine similarity between termination-guided directions constructed from pruned DeepSeek models on MATH-500. We compare only checkpoints with sufficient positive-terminated and negative non-terminating samples; for example, all 1.5B checkpoints have at least 50 samples in each class. Directions constructed from different pruned checkpoints exhibit high cosine similarity.

checkpoints A and B by the layer-averaged cosine similarity $\mathrm{sim}(A, B) = \frac{1}{L}\sum_{\ell=1}^{L} \cos(\mathbf{r}_\ell^A, \mathbf{r}_\ell^B)$ (details in Appendix D.5). For the analysis in this section we use MATH-500 traces; the directions used in all experiments (Section 5) are constructed from OpenR1-Math traces disjoint from every evaluation set.

4.2 Direction Consistency Across Checkpoints

Figure 5 reports $\mathrm{sim}(A, B)$ between directions from different pruning levels of the same model, and shows that they are highly consistent. For DS-Qwen-7B on MATH-500, the cosine similarity between directions from 50–70% pruning ratios remains at least 0.87; for DS-Qwen-1.5B, directions from 40–60% have pairwise cosine similarities of at least 0.83 (sample counts and results on OpenR1-Math are in Appendix C.4.5). These results suggest that the difference between repetitive reasoning and commitment is represented in a relatively stable direction in hidden-state space rather than being specific to an individual pruning checkpoint, enabling a single direction per model to generalize across sparsity levels. Appendix D.5 shows that this similarity is not an artifact of comparing checkpoints of the same model: across 15 checkpoint pairs, it exceeds both a label-shuffled baseline and a shared contrast unrelated to commitment in every pair.

4.3 Steering at Inference

TGAS constructs the direction at an intermediate sparsity level, where both $\mathcal{C}_b^+$ and $\mathcal{C}_b^-$ contain enough trajectories, and reuses it for all sparsity levels of the same model. We select the steering layer $\ell^\star$ on held-out branch points as the layer whose direction best separates the two classes (highest ROC AUC of the projection; Appendix C.4.5). The steering vector is this direction rescaled to the typical activation norm at that layer, $\mathbf{v} = \bar{n}_{\ell^\star}\, \mathbf{r}_{\ell^\star} / \|\mathbf{r}_{\ell^\star}\|_2$, where $\bar{n}_{\ell^\star}$ is the mean ℓ_2 norm of residual-stream representations over calibration trajectories. At layer $\ell^\star$, we add it to the residual stream at every token position t during autoregressive generation:

$$\tilde{\mathbf{h}}_{\ell^\star,t} \leftarrow \mathbf{h}_{\ell^\star,t} + \alpha\, \mathbf{v},$$

where the coefficient α scales the intervention relative to the typical activation norm; positive values steer toward commitment, while negative values steer in the opposite direction. We select α per checkpoint by accuracy on a held-out OpenR1-Math split, detailed in Table 12.

5 Experiments

Setting. We compare TGAS with three training-free post-pruning baselines: Forced-Answer, which stops after D reasoning tokens; No-Repeat, which suppresses repeated n-grams; and SEAL, which constructs a steering vector from reflection versus execution steps to reduce redundant reasoning. TGAS constructs one direction per model from OpenR1-Math calibration rollouts of an intermediate-sparsity checkpoint (50% for 1.5B models and 60% for other models), and reuses it across all sparsity levels. We also compare with training-based *Masked-LoRA* (Kim et al., 2024), which trains rank-64 LoRA adapters restricted to the surviving weights, details in Appendix C.4. Hyperparameters are tuned on a held-out OpenR1 split and evaluated on AIME-25, MATH-500, and LCB (Appendix C.4). Table 5 reports three representative models (1.5B–14B) at the two highest sparsity levels before accuracy collapses on AIME-25 and LCB; results for all six models are provided in Appendix E.1

Does Targeting Termination Recover Accuracy? Table 5 compares TGAS with training-free baselines. Consistent with our diagnosis, steering toward termination improves both termination and accuracy: averaged over benchmarks, TGAS improves over standard inference at every model-sparsity setting (by 1.3–3.7 points), and it raises termination in 16 of 18 checkpoint-benchmark pairs without lowering it in any. It also attains the best average accuracy among training-free methods in all six settings, although per-benchmark gains are often within one standard

Model	Ckpt	Method	MATH-500 Acc↑	MATH-500 Term↑	AIME-25 Acc↑	AIME-25 Term↑	LiveCodeBench Acc↑	LiveCodeBench Term↑	Avg Acc↑
OT3-1.5B	dense	Base	80.2±1.6	87.2±1.9	32.5±1.7	39.2±3.2	43.2±2.9	57.4±3.3	52.0
	obc_10	Base	79.6±1.3	87.0±0.5	34.2±8.3	40.0±2.7	43.4±2.3	57.1±2.6	52.4
		No-Repeat	79.2±1.9	86.6±2.2	35.0±6.9	41.7±6.9	38.9±2.6	55.5±1.6	51.0
		Forced-Answer	77.3±3.0	—	29.2±6.3	—	40.8±2.6	—	49.1
		SEAL	77.1±1.9	87.0±2.0	**41.7**±5.8	52.5±5.7	38.7±2.0	61.3±5.8	52.5
		TGAS	**81.7**±1.9	88.5±0.8	37.5±3.2	46.7±9.8	**45.5**±1.6	61.3±2.8	**54.9**$_{\uparrow 2.4}$
	obc_50	Base	56.9±3.8	67.4±4.0	**20.8**±1.7	32.5±5.0	10.0±0.6	22.4±1.3	29.2
		No-Repeat	54.4±0.3	67.4±1.3	16.7±4.7	25.8±3.2	6.3±0.9	23.2±4.5	25.8
		Forced-Answer	57.1±3.5	—	17.5±5.0	—	11.1±1.8	—	28.6
		SEAL	55.0±5.1	73.7±2.1	18.3±4.3	34.2±1.7	9.2±1.0	26.8±2.8	27.5
		TGAS	**58.4**±2.2	74.4±0.3	19.2±7.4	44.2±5.7	**13.9**±2.3	29.2±5.6	**30.5**$_{\uparrow 1.3}$
DS-Llama -8B	dense	Base	82.4±2.7	97.9±0.8	34.2±5.7	70.0±11.2	58.7±2.5	85.5±1.8	58.4
	obc_50	Base	65.5±0.3	88.5±2.2	21.7±7.9	75.0±8.4	42.6±0.6	66.8±1.8	43.3
		No-Repeat	64.7±0.8	90.6±1.3	20.8±8.8	76.7±2.7	33.2±3.0	75.3±4.9	39.6
		Forced-Answer	64.5±0.5	—	16.7±6.1	—	40.8±1.3	—	40.7
		SEAL	63.0±0.6	89.3±1.1	**23.3**±5.4	76.7±9.4	41.3±4.1	75.3±3.9	42.5
		TGAS	**67.6**±1.6	88.5±1.1	**23.3**±3.8	80.8±4.2	**43.9**±3.0	72.1±2.0	**44.9**$_{\uparrow 1.6}$
	obc_60	Base	30.0±3.0	50.6±2.4	3.3±2.7	29.2±4.2	10.8±1.6	26.3±2.3	14.7
		No-Repeat	33.4±0.8	74.4±1.1	2.5±3.2	50.8±10.7	9.5±1.5	42.1±3.7	15.1
		Forced-Answer	35.7±1.3	—	3.3±2.7	—	**11.1**±1.4	—	16.7
		SEAL	32.8±1.6	61.8±1.3	4.2±5.0	45.0±4.3	10.0±3.0	32.6±1.9	15.7
		TGAS	**36.6**±1.6	65.6±0.5	**7.5**±1.7	45.0±4.3	10.8±1.8	33.7±1.5	**18.3**$_{\uparrow 1.6}$
DS-Qwen -14B	dense	Base	88.4±0.3	99.0±0.3	54.2±5.0	87.5±5.0	73.2±3.9	90.8±2.2	71.9
	obc_50	Base	82.4±0.5	95.4±0.5	35.8±5.0	63.3±6.1	62.4±1.0	75.8±2.3	60.2
		No-Repeat	83.6±1.1	95.0±1.6	37.5±3.2	64.2±6.3	46.6±2.9	80.3±3.6	55.9
		Forced-Answer	81.1±1.9	—	36.7±4.7	—	56.6±3.3	—	58.1
		SEAL	84.4±0.8	98.1±0.5	**42.5**±5.7	70.0±3.8	61.3±2.8	78.7±2.9	62.7
		TGAS	**86.3**±0.5	95.4±0.3	**42.5**±3.2	70.8±1.7	**62.9**±3.5	77.1±3.5	**63.9**$_{\uparrow 1.2}$
	obc_60	Base	67.6±1.6	80.2±0.5	26.7±4.7	39.2±5.0	37.6±5.6	46.6±3.9	44.0
		No-Repeat	71.6±1.3	91.4±0.8	25.0±3.3	70.0±10.5	33.9±1.0	60.3±3.0	43.5
		Forced-Answer	70.8±0.8	—	**27.5**±3.2	—	38.9±4.3	—	45.7
		SEAL	71.0±0.8	86.6±1.8	24.2±4.2	45.0±6.4	**42.9**±2.9	60.0±0.9	46.0
		TGAS	**72.5**±1.3	88.9±2.2	25.8±3.2	44.2±4.2	**42.9**±2.6	53.7±3.1	**47.1**$_{\uparrow 1.1}$

Table 5: Training-free remedies for SparseGPT-pruned LRMs with OBC calibration. **Acc** is accuracy (%) and **Term** is the percentage of generations that terminate before the 16,384 token cap. For each checkpoint, the best accuracy is in **bold** and the second best is underlined. TGAS rows are shaded; the ↑ subscript in Avg denotes its gain over the best competitor.

deviation and TGAS trails the strongest baseline in four of 18 pairs. Baselines support our diagnosis: Forced-Answer, which also intervenes on stopping, is competitive when its budget matches problem difficulty; No-Repeat suppresses surface repetition but often hurts accuracy, especially on code; and SEAL, which suppresses reflection in general rather than looping specifically, raises termination at the cost of accuracy in several settings.

Does TGAS Fix the Diagnosed Failure? Table 7 shows that TGAS makes Δ_{cyc} more negative in all settings and substantially reduces decision points, roughly halving them at 60% sparsity ($12.3 \rightarrow 5.8$ and $11.5 \rightarrow 6.8$ per 1k tokens for 7B and 14B, detailed in Appendix F.1). TGAS also reduces generation length by 7.6% on average across all sparsity levels, including the dense model, and by 12.3% at 60% sparsity (Table 8). These results support our hypothesis that TGAS reduces repetitive reconsideration (Section 3.2).

Does TGAS Generalize Across Pruning Methods? On DS-Qwen-7B, Table 6 shows that TGAS consistently improves accuracy across pruning methods, outperforming the pruned baseline in all eight settings and achieving the best accuracy in seven. It also restores termination on degraded checkpoints, improving both accuracy and termination, unlike No-Repeat, which often trades accuracy for termination. At 60% sparsity, where RAC- and Wanda-pruned models retain little accuracy, all training-free methods remain far below the dense model (Appendix F.6, which also covers other models). Thus, TGAS transfers across pruning algorithms and calibration sets, provided the pruned model retains sufficient capability.

Comparison with Training-Based Recovery. Figure 6 shows that Masked-LoRA yields limited gains at 50% sparsity but substantially improves recovery at 60%, particularly on LCB, raising accuracy from 14.7% to 33.4% for DS-Qwen-7B and from 37.6% to 51.8% for 14B counterpart (Appendix E.2). This suggests that Masked-LoRA recovers capability lost to aggressive pruning. However, Masked-LoRA is trained on traces from the dense DS-Qwen-14B model, even for

(a) RAC (SparseGPT)

Ckpt	Method	AIME-25 Acc	AIME-25 Term	LiveCodeBench Acc	LiveCodeBench Term
rac_10	Base	35.8±5.0	79.2±4.2	56.3±1.8	85.0±3.8
	No-Repeat	38.3±5.8	80.0±5.4	47.4±3.1	85.8±1.1
	Forced-Answer	33.3±4.7	—	53.7±1.2	—
	TGAS	**39.2**±3.2	85.0±6.9	**58.7**±2.6	89.7±3.0
rac_50	Base	24.2±5.7	32.5±5.7	32.6±1.7	37.4±2.2
	No-Repeat	25.0±3.3	79.2±7.9	27.9±2.8	80.0±3.6
	Forced-Answer	23.3±3.8	—	**37.1**±1.8	—
	TGAS	**25.8**±1.7	59.2±7.4	34.5±4.0	47.9±1.1

(b) C4 calibration (Wanda)

Ckpt	Method	AIME-25 Acc	AIME-25 Term	LiveCodeBench Acc	LiveCodeBench Term
c4_40	Base	34.2±3.2	67.5±6.9	49.5±1.7	73.7±3.5
	No-Repeat	31.7±4.3	75.8±3.2	41.6±4.6	77.9±1.2
	Forced-Answer	30.0±4.7	—	49.7±2.0	—
	TGAS	**35.0**±1.9	75.8±5.0	**52.1**±0.6	81.6±4.7
c4_50	Base	19.2±1.7	47.5±9.2	31.8±5.3	52.9±5.6
	No-Repeat	15.8±3.2	61.7±1.9	26.6±4.7	67.4±2.4
	Forced-Answer	19.2±1.7	—	34.7±4.3	—
	TGAS	**21.7**±4.3	64.2±7.4	**39.2**±4.0	58.7±3.6

Table 6: Training-free remedies for checkpoints pruned with (a) RAC and (b) C4-calibrated Wanda on DS-Qwen-7B. **Acc** is accuracy (%) and **Term** the share of generations that terminated (%).

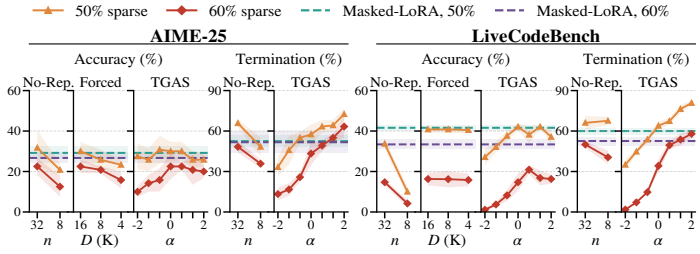

Figure 6: Ablation results on DS-Qwen-7B. TGAS provides controllable termination steering via α. Masked-LoRA exceeds training-free methods on LCB, but requires separate optimization for each model and sparsity level.

Sparsity	DS-Qwen 1.5B	DS-Qwen 7B	DS-Qwen 14B	DS-Llama 8B	OT3 1.5B	OT3 7B	Avg.	Saved
Dense	11.0	2.8	2.7	3.3	3.6	2.7	4.4	392
10%	10.2	10.9	5.9	4.6	5.9	3.7	6.9	517
50%	17.4	8.3	1.7	2.5	7.9	3.4	6.9	680
60%	23.8	18.5	9.4	12.9	3.6	5.7	**12.3**	**1,423**

Table 8: Reduction in generated tokens (%) from TGAS relative to the unsteered model, averaged over MATH-500, AIME-25, and LiveCodeBench. **Avg.**: mean reduction over models (%); **Saved**: mean tokens saved per generation.

Model	Ckpt	Δ_{cyc} at decisions U	Δ_{cyc} at decisions S	DPs/1k U	DPs/1k S
DS-Qwen -7B	obc_50	−15.3±0.5	−17.1±0.5	9.0	5.4
	obc_60	−12.0±0.4	−15.4±0.5	12.3	5.8
DS-Qwen -14B	obc_50	−14.7±0.4	−17.9±0.5	7.2	5.5
	obc_60	−12.7±0.4	−14.0±0.4	11.5	6.8

Table 7: Δ_{cyc} at decision points, defined in Section 3.2, on MATH-500, unsteered (**U**) vs. steered with TGAS (**S**). **DPs/1k**: decision points per 1,000 tokens.

smaller models, confounding recovery with distillation. This contrast supports our diagnosis: under moderate pruning, much of the loss stems from failure to stop, which TGAS addresses without training; at higher sparsity, some capability is genuinely lost and can only be restored by updating the weights. Such recovery, however, requires a separate adapter and optimization procedure for each sparsity level and model.

Ablation Studies. Figure 6 summarizes the ablations on DS-Qwen-7B (other models in Appendix F). For Forced-Answer, reducing the reasoning budget consistently hurts accuracy, with 4K performing worse than 16K on both AIME-25 and LCB. For No-Repeat, 8-gram blocking is substantially more harmful than 32-gram blocking, particularly at 60% sparsity. For TGAS, increasing α consistently improves termination, especially on LCB, so α provides a controllable knob for balancing termination and accuracy. We also test TGAS robustness to the number of samples and checkpoint choice, finding that 50 samples per class suffice (Appendix F.5) and directions from different checkpoints of the same model remain effective (Appendix F.7). We also compare TGAS with Plan-CoT (Sheng et al., 2025) and report significance tests with $K = 16$ on AIME-25 and LCB in Appendix F.9 and F.10, respectively.

6 Conclusion

We find that under moderate pruning, a substantial and recoverable share of accuracy loss stems from disrupted termination rather than incorrect reasoning. As sparsity increases, pruned LRMs exhibit reasoning divergence: they often reach the correct answer and recognize where to stop under teacher forcing, yet fall into repetitive loops and fail to commit within the token budget. The hidden-state contrast between terminating and looping continuations is consistent across pruned checkpoints; building on it, TGAS steers pruned models toward termination without retraining, improving termination and accuracy across configurations. At higher sparsity, some capability is genuinely lost and can only be recovered by updating the weights, suggesting that future pruning methods should preserve termination behavior.

Limitations

Our work focuses on post-pruning recovery and does not modify the pruning process itself. Thus, the proposed interventions recover degraded behavior after pruning rather than preventing it during pruning. Future work could incorporate these observations into pruning algorithms to preserve termination and reasoning behavior directly, alongside a more theoretical characterization of the relationship between pruning and reasoning dynamics.

Following prior work on pruning LRMs (Lucas et al., 2026; Nguyen et al., 2026; Ji et al., 2025), we study unstructured sparsity; semi-structured patterns, which SparseGPT and Wanda also support, are a natural extension. All models are distilled or SFT-trained LRMs of at most 14B parameters, and the benchmarks cover mathematics and code only; whether the diagnosis holds for RL-trained reasoning models, larger models, or open-ended tasks remains open. Our analysis uses a 16,384-token budget, and TGAS requires sampling continuations from a pruned checkpoint once per model to construct the direction. Finally, TGAS targets termination; at high sparsity, where capability is genuinely lost, it recovers only part of the gap to the dense model.

Ethical Considerations

This work studies inference-time interventions for pruned reasoning models and does not introduce new data collection involving human subjects. The proposed methods can alter model generation behavior, including its tendency to terminate or continue reasoning; therefore, they should be evaluated for task-specific reliability before deployment.

References

Andy Arditi, Oscar Obeso, Aaquib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. 2024. Refusal in language models is mediated by a single direction. *Advances in Neural Information Processing Systems*, 37:136037–136083.

Abhinav Bandari, Lu Yin, Cheng-Yu Hsieh, Ajay Kumar Jaiswal, Tianlong Chen, Li Shen, Ranjay Krishna, and Shiwei Liu. 2024. Is C4 dataset optimal for pruning? an investigation of calibration data for LLM pruning. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 18089–18099.

Qiguang Chen, Libo Qin, Jinhao Liu, Dengyun Peng, Jiannan Guan, Peng Wang, Mengkang Hu, Yuhang Zhou, Te Gao, and Wanxiang Che. 2026. Towards reasoning era: A survey of long chain-of-thought for reasoning large language models. *Science China Information Sciences*, 69(6):161101.

Runjin Chen, Zhenyu Zhang, Junyuan Hong, Souvik Kundu, and Zhangyang Wang. 2025. SEAL: Steerable reasoning calibration of large language models for free. In *Second Conference on Language Modeling*.

Hongrong Cheng, Miao Zhang, and Javen Qinfeng Shi. 2024. A survey on deep neural network pruning: Taxonomy, comparison, analysis, and recommendations. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(12):10558–10578.

Frances Ding, Jean-Stanislas Denain, and Jacob Steinhardt. 2021. Grounding representation similarity through statistical testing. *Advances in neural information processing systems*, 34:1556–1568.

Elias Frantar and Dan Alistarh. 2023. SparseGPT: Massive language models can be accurately pruned in one-shot. In *International conference on machine learning*, pages 10323–10337. PMLR.

Elias Frantar, Sidak Pal Singh, and Dan Alistarh. 2022. Optimal brain compression: A framework for accurate post-training quantization and pruning. *Advances in Neural Information Processing Systems*, 35:4475–4488.

Etash Guha, Ryan Marten, Sedrick Keh, Negin Raoof, Georgios Smyrnis, Hritik Bansal, Marianna Nezhurina, Jean Mercat, Trung Vu, Zayne Sprague, et al. 2026. OpenThoughts: Data recipes for reasoning models. In *International Conference on Learning Representations*, volume 2026, pages 108059–108130.

Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Peiyi Wang, Qihao Zhu, Runxin Xu, Ruoyu Zhang, Shirong Ma, Xiao Bi, et al. 2025. DeepSeek-R1: Incentivizing reasoning capability in LLMs via reinforcement learning. *arXiv preprint arXiv:2501.12948*.

Nathan Habib, Clémentine Fourrier, Hynek Kydlíček, Thomas Wolf, and Lewis Tunstall. 2023. Lighteval: A lightweight framework for llm evaluation.

Babak Hassibi and David Stork. 1992. Second order derivatives for network pruning: Optimal brain surgeon. In *Advances in Neural Information Processing Systems*, volume 5. Morgan-Kaufmann.

Bowei He, Lihao Yin, Huiling Zhen, Xiaokun Zhang, Mingxuan Yuan, and Chen Ma. 2026. Paser: Post-training data selection for efficient pruned large language model recovery. In *International Conference on Learning Representations*, volume 2026, pages 14510–14554.

Hugging Face. 2025. Open r1: A fully open reproduction of deepseek-r1.

Naman Jain, Alex Gu, Wen-Ding Li, Fanjia Yan, Tianjun Zhang, Sida Wang, Armando Solar-Lezama, Koushik Sen, and Ion Stoica. 2025. LiveCodeBench: Holistic and contamination free evaluation of large language models for code. In *International Conference on Learning Representations*, volume 2025, pages 58791–58831.

Yixin Ji, Yang Xiang, Juntao Li, Qingrong Xia, Ping Li, Xinyu Duan, Zhefeng Wang, and Min Zhang. 2025. Beware of calibration data for pruning large language models. In *International Conference on Learning Representations*, volume 2025, pages 18604–18620.

Bo-Kyeong Kim, Geonmin Kim, Tae-Ho Kim, Thibault Castells, Shinkook Choi, Junho Shin, and Hyoung-Kyu Song. 2024. Shortened LLaMA: Depth pruning for large language models with comparison of retraining methods. *arXiv preprint arXiv:2402.02834*.

Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph Gonzalez, Hao Zhang, and Ion Stoica. 2023. Efficient memory management for large language model serving with PagedAttention. In *Proceedings of the 29th symposium on operating systems principles*, pages 611–626.

Yann LeCun, John Denker, and Sara Solla. 1989. Optimal brain damage. In *Advances in Neural Information Processing Systems*, volume 2. Morgan-Kaufmann.

Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. 2024. Let's verify step by step. In *International Conference on Learning Representations*, volume 2024, pages 39578–39601.

Ruikang Liu, Yuxuan Sun, Manyi Zhang, Haoli Bai, Xianzhi Yu, Tiezheng YU, Chun Yuan, and Lu Hou. 2025a. Quantization hurts reasoning? an empirical study on quantized reasoning models. In *Second Conference on Language Modeling*.

Yijiang Liu, Huanrui Yang, Youxin Chen, Rongyu Zhang, Miao Wang, Yuan Du, and Li Du. 2025b. PAT: Pruning-aware tuning for large language models. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 24686–24695.

Sanae Lotfi, Polina Kirichenko, Steven Li, and Zechun Liu. 2026. Quantized reasoning models think they need to think longer, but they do not. *arXiv preprint arXiv:2606.00206*.

Ryan Lucas, Kayhan Behdin, Zhipeng Wang, Qingquan Song, Shao Tang, and Rahul Mazumder. 2026. Reasoning models can be accurately pruned via chain-of-thought reconstruction. In *International Conference on Learning Representations*.

Xinyin Ma, Gongfan Fang, and Xinchao Wang. 2023. LLM-Pruner: On the structural pruning of large language models. *Advances in neural information processing systems*, 36:21702–21720.

Xin Men, Mingyu Xu, Qingyu Zhang, Qianhao Yuan, Bingning Wang, Hongyu Lin, Yaojie Lu, Xianpei Han, and Weipeng Chen. 2025. ShortGPT: Layers in large language models are more redundant than you expect. In *Findings of the Association for Computational Linguistics: ACL 2025*, pages 20192–20204.

Saurav Muralidharan, Sharath Turuvekere Sreenivas, Raviraj Joshi, Marcin Chochowski, Mostofa Patwary, Mohammad Shoeybi, Bryan Catanzaro, Jan Kautz, and Pavlo Molchanov. 2024. Compact language models via pruning and knowledge distillation. *Advances in Neural Information Processing Systems*, 37:41076–41102.

Ha Lan Nguyen, Huy Hoang Tran, Trac-Duy Tran, and Dung D. Le. 2026. OBC-Prune: Outcome-based calibration for large reasoning model pruning.

Charilaos Pipis, Shivam Garg, Vasilis Kontonis, Vaishnavi Shrivastava, Akshay Krishnamurthy, and Dimitris Papailiopoulos. 2026. Wait, wait, wait... why do reasoning models loop? In *International Conference on Machine Learning*.

Leheng Sheng, An Zhang, Zijian Wu, Weixiang Zhao, Changshuo Shen, Yi Zhang, Xiang Wang, and Tat-Seng Chua. 2025. On reasoning strength planning in large reasoning models. In *Advances in Neural Information Processing Systems*, volume 38, pages 62978–63012. Curran Associates, Inc.

Mingjie Sun, Zhuang Liu, Anna Bair, and J. Zico Kolter. 2024. A simple and effective pruning approach for large language models. In *International Conference on Learning Representations*.

Qwen Team. 2026. Qwen3.5: Towards native multimodal agents.

Ziyan Wang, Enmao Diao, Qi Le, Pu Wang, Guanchu Wang, Minwoo Lee, Shu-ping Yeh, and Li Yang. 2025. Think before you prune: Self-reflective structured pruning for reasoning language models. *arXiv preprint arXiv:2512.02185*.

Miles Williams and Nikolaos Aletras. 2024. On the impact of calibration data in post-training quantization and pruning. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 10100–10118.

Mengzhou Xia, Tianyu Gao, Zhiyuan Zeng, and Danqi Chen. 2024a. Sheared LLaMA: Accelerating language model pre-training via structured pruning. In *International Conference on Learning Representations*, volume 2024, pages 5385–5409.

Mengzhou Xia, Sadhika Malladi, Suchin Gururangan, Sanjeev Arora, and Danqi Chen. 2024b. Less: Selecting influential data for targeted instruction tuning.

Yang Xiang, Yixin Ji, Juntao Li, and Min Zhang. 2025. Think before you prune: Selective self-generated calibration for pruning large reasoning models. *arXiv preprint arXiv:2511.18864*.

Duzhen Zhang, Zhong-Zhi Li, Ming-Liang Zhang, Jiaxin Zhang, Zengyan Liu, Yuxuan Yao, Haotian Xu, Junhao Zheng, Xiuyi Chen, Yingying Zhang, et al. 2026a. From system 1 to system 2: A survey of reasoning large language models. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 48(3):3335–3354.

Nan Zhang, Eugene Kwek, Yusen Zhang, Ngoc-Hieu Nguyen, Prasenjit Mitra, and Rui Zhang. 2026b. When reasoning meets compression: Understanding the effects of LLMs compression on large reasoning models. In *The Fourteenth International Conference on Learning Representations*.

Yifan Zhang and Team Math-AI. 2025. American invitational mathematics examination (AIME) 2025.

Appendices

A Related Works **11**

A.1 Pruning Methods for Large Language Models 11

A.2 Pruning Methods for Large Reasoning Models 12

B Preliminaries **12**

C Detailed Experimental Setup **13**

C.1 Generation Setup 13

C.2 Measuring Answer Presence in Reasoning Divergence 14

C.3 Independent Validation with an LLM Judge 14

C.4 Training Baseline 15

D Complete Observation Results **20**

D.1 Answer Retention and Termination in Pruned Models on Dense-Model-Correct Samples 20

D.2 Comprehensive Evaluation of Larger Budget Generation 20

D.3 Comprehensive Evaluation of the Stop-Token Mechanism 20

D.4 Comprehensive Evaluation of Transition of Pruned LRMs 21

D.5 Comprehensive Commitment Directions Analysis 22

E Additional Main Experimental Results **27**

E.1 Comprehensive Evaluation of Training-free Methods 27

E.2 Comparison of Training-Free and Training-Based Methods 27

E.3 Headroom and Recovery Analysis 27

F Additional Ablation Studies **32**

F.1 Behavioral Changes in Pruned Models After Steering 32

F.2 Ablation Study of Forced-Answer 32

F.3 Ablation Study of No-Repeat . . . 32

F.4 Ablation Study of Steering Strength 32

F.5 Ablation on Number of Sampled Created Directions 33

F.6 Generalization Across Pruning Methods 34

F.7 Cross-Testing with Different Directions 35

F.8 Comparison with SEAL 35

F.9 Comparison with Plan-CoT 35

F.10 Statistical Significance of TGAS Gains 36

A Related Works

A.1 Pruning Methods for Large Language Models

Pruning has long been explored as a means of reducing the computational and memory costs of neural networks while preserving model performance (LeCun et al., 1989; Hassibi and Stork, 1992). Recent work has extended pruning to large language models (LLMs), with methods broadly falling into unstructured weight pruning and structured architectural pruning (Cheng et al., 2024). For unstructured pruning, Frantar and Alistarh (2023) proposed SparseGPT, a post-training pruning method based on approximate second-order information that enables one-shot pruning of GPT-scale models without subsequent retraining. In contrast, Wanda (Sun et al., 2024) estimates weight importance from the product of

weight magnitude and input activation magnitude, providing a simple activation-aware pruning criterion without weight reconstruction or retraining.

For structured pruning, LLM-Pruner (Ma et al., 2023) removes coupled structures according to gradient-based importance and subsequently recovers performance through lightweight fine-tuning. ShortGPT (Men et al., 2025) instead identifies redundant transformer blocks using block influence and directly removes entire layers. Sheared-LLaMA (Xia et al., 2024a) further considers pruning layers and attention heads, followed by continued pre-training to recover the capacity of the compressed model. Kim et al. (2024) further shows that simple depth pruning can effectively compress LLMs while achieving comparable or superior performance to other post-pruning fine-tuning methods. Complementary work examines calibration data in post-training pruning (Bandari et al., 2024; Williams and Aletras, 2024). While pruning often uses small calibration sets, Ji et al. (2025) show that calibration choice substantially affects pruning quality, especially at high sparsity, and that self-generated data can be an effective alternative when pre-training data are unavailable.

A.2 Pruning Methods for Large Reasoning Models

Recent LRMs, such as DeepSeek-R1 (Guo et al., 2025) and Qwen3.5 (Team, 2026), introduce a distinct challenge for model compression. Unlike conventional LLMs, whose computation is often dominated by processing the input context, LRMs can generate long chain-of-thought (CoT) trajectories during autoregressive decoding. Consequently, preserving the input-to-output behavior alone may not be sufficient to preserve reasoning capability (Chen et al., 2026). Recent studies have therefore begun to adapt pruning to the decode-dominated nature of LRMs (He et al., 2026; Xia et al., 2024b). Lucas et al. (2026) show that conventional pruning can degrade reasoning performance and increase inference cost by inducing longer reasoning trajectories. They attribute this to the mismatch between input-oriented activation reconstruction and decode-time reasoning behavior, and propose Reasoning-Aware Compression (RAC), which incorporates both input tokens and on-policy reasoning trajectories.

Complementing this direction, Wang et al. (2025) propose RESP, a self-reflective structured pruning framework that explicitly aligns pruning decisions with the reasoning dynamics of reasoning language models. Their approach uses self-generated reasoning traces as calibration data, decode-only gradient-based importance estimation, and progressive regeneration of calibration traces during pruning. Building on these findings, Xiang et al. (2025) argue that model-generated reasoning trajectories offer a more informative calibration signal for pruning LRMs than conventional human-written or generic language-modeling data.

Beyond individual pruning methods, Zhang et al. (2026b) systematically investigate how different compression techniques, including pruning, quantization, and distillation, affect large reasoning models, with a particular focus on the components associated with reasoning degradation. Their findings suggest that compression can have distinct effects on memorization and reasoning, and that certain components, particularly MLP modules in later layers, play a disproportionately important role in preserving reasoning capabilities. Complementary studies on quantization further show that reasoning models exhibit distinct sensitivity to reduced numerical precision compared with conventional LLMs (Liu et al., 2025a). More recently, Lotfi et al. (2026) investigate the overthinking phenomenon in quantized reasoning models, highlighting another potential consequence of model compression on the reasoning behavior of LRMs.

B Preliminaries

Pruning reduces an LLM's size by setting selected weights to zero while preserving its performance. Most one-shot methods formulate pruning as a layer-wise reconstruction problem (Frantar et al., 2022), finding a sparse weight matrix $\hat{\mathbf{W}}_\ell$ that preserves the original layer outputs on representative calibration inputs $\hat{\mathbf{W}}_\ell = \operatorname{argmin}_{\|\hat{\mathbf{W}}_\ell\|_0 \le S} \|\mathbf{W}_\ell \mathbf{X}_\ell - \hat{\mathbf{W}}_\ell \mathbf{X}_\ell\|_F^2$. guides which weights can be removed, while S specifies how many weights are retained.

In this work, we focus on two representative unstructured one-shot pruning methods, SparseGPT (Frantar and Alistarh, 2023) and Wanda (Sun et al., 2024), due to their computational efficiency and ability to prune large language models without costly retraining. SparseGPT (Frantar and Alistarh, 2023) scores weights via the layer Hessian

$\mathbf{H}_\ell = \mathbf{X}_\ell \mathbf{X}_\ell^\top$, using $w_q^2/(\mathbf{H}_\ell^{-1})_{qq}$ and compensating for removal-induced errors. Wanda (Sun et al., 2024) instead uses the diagonal score $s_{ij} = |w_{ij}| \cdot \|\mathbf{x}_j\|_2$ without compensation.

Both methods depend on calibration activations, but standard text calibration may not adequately capture the activation distribution during extended reasoning (Liu et al., 2025a). We therefore combine one-shot pruning methods with RAC (Lucas et al., 2026) and OBC (Nguyen et al., 2026) to construct calibration sets tailored to pruned LRMs. Specifically, RAC augments prompt activations with model-generated chain-of-thought activations, forming $\mathbf{X}_\ell = [\mathbf{X}_\ell^{\text{prompt}}, \mathbf{X}_\ell^{\text{CoT}}]$ to estimate parameter importance using inputs that better reflect deployment-time reasoning. Besides, OBC generates difficulty-matched correct and incorrect rollouts and assigns token-level calibration weights according to each reasoning step's causal importance. These weights rescale calibration activations before one-shot pruning, helping preserve parameters critical to correct reasoning.

C Detailed Experimental Setup

This appendix describes the experimental protocol used to investigate the behavioral changes induced by weight pruning.

C.1 Generation Setup

Setting. We evaluate DeepSeek-R1-Distill-Qwen, DeepSeek-R1-Distill-Llama (Guo et al., 2025), and OpenThinker3 (Guha et al., 2026) under one-shot unstructured pruning of decoder linear layers at varying sparsity levels. We construct calibration data using OBC, RAC, or C4 (Bandari et al., 2024), with SparseGPT and Wanda as pruning backbones. We consider one-shot unstructured pruning at nominal sparsities of 10-70%. Pruning is restricted to decoder linear layers, including the MLP projections, full-attention $q/k/v/o$ projections. Embeddings and `lm_head` remain dense, followed Frantar and Alistarh (2023).

- **OBC/SparseGPT**: OBC constructs difficulty-matched correct and incorrect rollouts and assigns per-token calibration weights based on the causal importance of reasoning steps. These weights rescale calibration activations before one-shot pruning methods, preserving parameters that are more important for correct reasoning.
- **RAC/SparseGPT**: The dense LRM generates chain-of-thought trajectories for the calibration prompts, and hidden states collected during autoregressive decoding are concatenated with prompt activations before being passed to SparseGPT. This produces a calibration distribution that better reflects the decode-heavy reasoning workload.
- **C4/Wanda**: the same C4 calibration set and pruning plan, using Wanda's activation-norm criterion instead of SparseGPT's second-order statistic.

Our evaluation covers 262 Level 4–5 MATH-500 questions, 30 AIME-25 questions, and 95 easy/medium LiveCodeBench questions. We evaluate on held-out problems from three benchmarks spanning mathematical reasoning, competition-level problems, and code generation:

- **MATH-500**: 262 samples, covering two most difficult levels and seven subjects.
- **AIME-25**: all 30 problems, providing a challenging multi-step mathematical reasoning benchmark.
- **LiveCodeBench-v6**: 95 easy and medium problems from code contests held in 2025. These problems post-date the models' training data and therefore provide a test of generalization rather than memorization. Since steering vectors are constructed from mathematical trajectories, LiveCodeBench additionally evaluates cross-domain transfer.

Unless otherwise specified, generation uses a 16,384-token budget, $T = 0.6$ and top-$p = 0.95$. For each question, we generate $k = 4$ responses for Math-500 and LCB, and AIME-25, reporting the average for more stable evaluation. We use vLLM (Kwon et al., 2023) for inference and Lighteval (Habib et al., 2023) for evaluation.

Computing infrastructure. All experiments were run on 1 NVIDIA H200 144GB GPU. Pruning a checkpoint with SparseGPT/Wanda takes about 1 GPU hour; constructing a TGAS direction (sampling continuations and computing hidden

states) takes about 1 GPU hour per model; evaluation across all benchmarks takes about 2 GPU hours per checkpoint and method. Masked-LoRA training takes about 6 GPU hours per checkpoint.

C.2 Measuring Answer Presence in Reasoning Divergence

Final-answer correctness alone does not determine whether a model has successfully reasoned to the correct solution. In particular, a non-termination trace may contain the correct answer in its reasoning segment even though the generation never reaches the final-answer stage. We therefore introduce a separate *answer-presence* measurement that asks whether the reasoning process itself reaches or asserts the gold answer.

We therefore modify the extraction procedure to identify candidate values asserted throughout the reasoning trace, rather than extracting a single final answer, and evaluate each candidate using the same grading procedure. To account for the many intermediate numerical values in long traces, we define a hierarchy of increasingly permissive answer-presence criteria:

- **Explicit**: an explicit answer in the think segment, identified by either a `\boxed{}` or `\fbox{}` expression matching the gold answer, or a statement of the form *"the (final) answer is X"* or *"Answer: X"*, where the remainder of the line is taken as X.
- **Conclusive**: an *Explicit* answer or an *Assertion*, where an *Assertion* is a sentence containing a conclusion cue, detailed *"so", "thus", "therefore", "hence", "the answer", "which is/gives/means/equals", "we get", or "in total")* whose text after the last copula with *"is", "equals", "gives", "yields", "="*, or *"≈"* matches the gold answer, following Lotfi et al. (2026).

The nested design distinguishes explicit commitment from incidental occurrences. In particular, *Conclusive* is our primary answer-presence criterion, as it captures cases where the model reaches the correct result but fails to transition to a final answer. Figure 7 shows that many incorrectly graded traces nevertheless contain the correct answer with Explicit or Conclusive mentioned. At 60% sparsity, this occurs in 41.4% and 39.2% of non-terminating traces for DeepSeek-R1-Distill-Qwen 7B and 14B, respectively. These results show that correct answers are often already present in non-termination reasoning traces, rather than appearing only as incidental intermediate values.

Table 9: Agreement between our automated measurements and the independent Gemini-2.5-Flash judge.

comparison	n	agreement	κ
Conclusive vs judge	250	96.0%	0.898
Explicit vs judge	250	84.8%	0.782

C.3 Independent Validation with an LLM Judge

To independently validate both our deterministic answer-grading procedure and the answer-presence measurement, we use Gemini-2.5-Flash[2] as an external judge on the responses generated by the DeepSeek-R1-Distill-Qwen-7B model. The judge is used strictly for validation: its outputs are never written back to the dataset, used to modify the correct labels, or included in the main analysis. We validate answer presence in the reasoning trace with detailed prompts in Figure 8.

LLM-as-a-judge setup. For independent validation, we use Gemini-2.5-Flash at temperature 0. The judge evaluates answer-presence agreement whether the reasoning trace explicitly asserts the gold result. We sample 50 traces from each of five sparsity levels ($10\%, 40\% - 70\%$), 250 in total.

Answer-presence validation. Given the problem, gold answer, and full reasoning trace, the judge determines whether the model explicitly arrives at the gold answer, even if the trace does not terminate. The judge must not solve the problem independently or infer an answer from incomplete reasoning. It must also distinguish the final answer from intermediate calculations, rejected candidates, copied values, and coincidental matches.

Agreement with the independent judge. The independent evaluation shows strong agreement with our automated measurements. On the same random sample, *Conclusive* achieves 96.0% agreement with the judge and $\kappa = 0.898$, while the more conservative *Explicit* criterion achieves 84.8% agreement and $\kappa = 0.782$. These results

[2] ai.google.dev/gemini-api/docs/models/gemini-2.5-flash

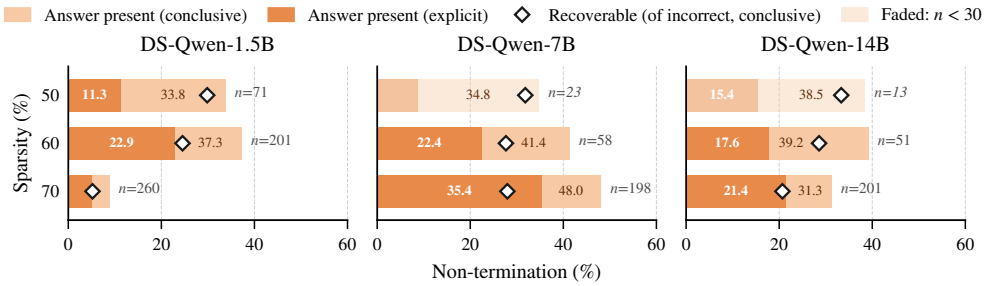

Figure 7: Gold answers appearing in reasoning traces graded incorrect for DeepSeek-R1-Distill-Qwen models. *Explicit* and *Conclusive* denote gold-answer presence under nested detection rules.

ANSWER_PRESENT_PROMPT

```
You are analysing a language model's chain-of-thought for a math problem. The model often runs out of
tokens before finishing, so it may never state a final answer.

PROBLEM:
{problem}

CORRECT ANSWER: {gold}

MODEL'S REASONING TRACE:
<{think}>

Question: does the reasoning trace state {gold} AS ITS OWN RESULT
for the problem?

Answer TRUE only if the trace asserts that value as the answer or as the result of its computation,
for example "so the answer is X", "\boxed{X}", "therefore the angle is X",
or a concluding line that evaluates to X.

Answer FALSE if the value only appears as:
  - a quantity given in the problem statement, or restated from it,
  - an intermediate step that is not presented as the answer,
  - a candidate the model considers and then rejects,
  - a coincidental appearance of the same number in unrelated arithmetic.

Judge only what the trace actually says. Do not solve the problem yourself, and do not credit
the trace for reasoning that would have led to the answer  if continued.

Respond with JSON only: {"present": true or false, "quote": "<short verbatim quote that states it,
or empty>", "reason": "<one short sentence>"}
```

Figure 8: Prompt used to independently validate whether a reasoning trace asserts the reference answer as its own result.

support *Conclusive* as the primary measure of answer presence and confirm the reliability of our automated evaluation, detailed in Table 9.

C.4 Training Baseline

C.4.1 Masked LoRA Recovery

To assess whether pruning degradation can be recovered without restoring zeroed weights, we extend masked fine-tuning (Ma et al., 2023; Xia et al., 2024a; Kim et al., 2024) to LRMs by restricting LoRA updates strictly to the surviving support. For a pruned weight $\widehat{W}_\ell$ and its binary support Ω_ℓ, the update is defined as:

$$\widetilde{W}_\ell = \widehat{W}_\ell + s \cdot \big(P_\ell Q_\ell \odot \Omega_\ell\big), \qquad [\Omega_\ell]_{ij} = \mathbf{1}\left[[\widehat{W}_\ell]_{ij} \neq 0\right] \tag{2}$$

where P_ℓ and Q_ℓ are trainable low-rank adapters with rank $\rho = 64$, and Ω_ℓ masks pruned positions to preserve sparsity. We apply LoRA to 200 decoder linear layers, excluding gate matrices, embeddings, and tied heads, and train for one epoch. After training, we merge the adapters and verify sparsity. We evaluate accuracy and termination rate to assess recovery from non-termination.

Supervision signal. We train on verified-correct reasoning traces from OpenR1-Math, generated by dense DeepSeek-R1-Distill-Qwen 14B. Target traces are fixed across all checkpoints so that every model is trained on identical supervision. Prompts are masked, supervising only the response span $\mathcal{T}(x)$ over a gradient-accumulation group $\mathcal{G}$:

$$\mathcal{L}_{\mathcal{G}}(\theta) = -\frac{1}{N_{\mathcal{G}}} \sum_{x \in \mathcal{G}} \sum_{t \in \mathcal{T}(x)} \log p_\theta\big(x_{t+1} \mid x_{\le t}\big),$$
$$N_{\mathcal{G}} = \sum_{x \in \mathcal{G}} |\mathcal{T}(x)|\,.$$

Normalizing by the total number of supervised tokens avoids bias toward shorter sequences. Since Masked-LoRA uses verified reasoning traces from a dense 14B model, its performance reflects both recovery and distillation, particularly for smaller models. We therefore treat it as a reference for training-free methods, serving as an upper bound only at high sparsity (60%–70% sparse ratios).

C.4.2 Forced-Termination Intervention.

We investigate whether non-termination in pruned models can be mitigated by explicitly enforcing termination after sufficiently long reasoning traces. Our analysis shows that pruned models largely preserve both reasoning traces and the stopping mechanism of dense models, particularly at moderate sparsity, motivating this intervention. We therefore introduce a simple forced-termination intervention: given a reasoning trace, we append the `</think>` token to terminate the reasoning phase and then prompt the model to produce the final answer using an explicit pattern such as "Answer:" after D generation budgets. This intervention prevents further free-running generation and directly tests whether the existing reasoning trace is sufficient to recover the final answer without modifying the model parameters.

In practice, the difficulty of a user query is unknown, making it difficult to select an appropriate D. Moreover, tuning D on a development set may not generalize across benchmarks with different difficulty distributions. We therefore use a fixed budget of $D = 16{,}384$ tokens for AIME-25 and $D = 8,192$ for MATH-500 and LCB and evaluate alternative budgets in Appendix F.2.

C.4.3 No-Repeat

Our observations in Section 3 show that unconstrained decoding of pruned LRMs can fall into cyclic loops. To prevent this, we forbid the model from repeating any n-gram already present in its context. Let $y_{<t} = (y_1, \ldots, y_{t-1})$ be the tokens generated so far and $\mathbf{z}_t \in \mathbb{R}^{|\mathcal{V}|}$ the logits at step t. We define the set of banned tokens as those that would complete an n-gram whose $(n-1)$-token prefix matches the current suffix: $\mathcal{B}_t = \{y_{i+n-1} \mid 1 \le i \le t-n,\ (y_i, \ldots, y_{i+n-2}) = (y_{t-n+1}, \ldots, y_{t-1})\}$. The adjusted logits are:

$$\tilde{z}_{t,v} = \begin{cases} -\infty, & v \in \mathcal{B}_t, \\ z_{t,v}, & \text{otherwise}, \end{cases} \tag{3}$$

and the next token is drawn from $\mathrm{softmax}(\tilde{\mathbf{z}}_t)$[3].

We evaluate n-gram sizes in $\{8, 32\}$ on the held-out OpenR1-Math set and select $n = 32$ as the default based on its performance. Ablation results are reported in Appendix F.3.

C.4.4 SEAL: Steerable Reasoning Calibration of LRMs

SEAL (Chen et al., 2025) is a training-free activation-steering method that reduces redundant reasoning in LRMs. It segments a reasoning trace into thoughts separated by `\n\n` and labels each thought by keyword rules: a thought is a *transition* if it starts with *"Alternatively"* or contains phrases such as *"another approach"*, a *reflection* if it starts with *"Wait"* or contains phrases such as *"let me check"* or *"make sure"*, and an *execution* thought otherwise. SEAL finds that, for problems of comparable difficulty, incorrect responses contain more reflection and transition thoughts than correct ones. Since these thought types are separable in the representation space, SEAL constructs activation steering vectors to encourage desired reasoning behaviors, steering LRMs toward more *execution* thoughts.

To extract a steering vector, SEAL collects the layer-ℓ hidden states at the `\n\n` token closing each thought on a calibration set, averages them per category, and takes the difference between execution thoughts and the union of reflection and transition thoughts.

$$\mathbf{S} = \bar{\mathbf{H}}_E - \bar{\mathbf{H}}_{RT},$$

with E, R, and T denote *execution*, *reflection*, and *transition* thoughts, respectively. During decoding, SEAL adds $\alpha\,\mathbf{S}$ to the hidden state of every `\n\n` token at layer ℓ, so the intervention is applied only at step boundaries. We construct $\mathbf{S}$ from calibration rollouts on OpenR1-Math, the

[3] Implemented with `no_repeat_ngram_size` in Hugging Face `transformers`.

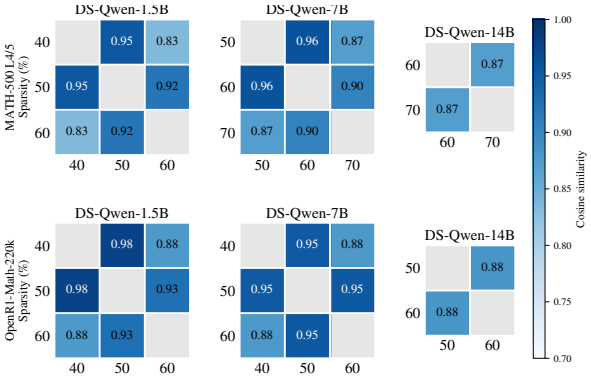

Figure 9: Cosine similarity between termination-guided directions from DeepSeek-R1-Distill-Qwen checkpoints. We include only checkpoints with sufficient positive and negative samples.

same source used for TGAS, apply it at the layers listed in Table 12, and select its strength from $\{0.25, 0.5, \ldots, 1.25, 1.5\}$ per checkpoint.

C.4.5 Termination-guided Activation Steering

Figure 9 provides an additional observation that directions constructed from different pruned checkpoints can remain highly similar, suggesting a shared behavioral direction across pruning levels. A detailed analysis is provided in Appendix D.5.

To test whether this direction generalizes beyond the calibration data, our benchmarking uses an out-of-distribution (OOD) setting: we construct one steering vector per model, from its medium-sparsity checkpoint and calibration rollouts on `open-r1/OpenR1-Math`, using at most six rollouts per problem per class, and evaluate it on separate downstream benchmarks. We evaluate on MATH-500, AIME-25, and LiveCodeBench-v6 to test generalization across unseen mathematical problems and a different domain. Each vector is constructed once and fixed across all pruning conditions, with calibration problems disjoint from evaluation sets. We provide the statistic of number of positive and negative continuations in Table 10 and Table 11.

To select the steering layer, we split the branch points $\mathcal{B}$ into a fitting set $\mathcal{B}_{\text{fit}}$ and a held-out set $\mathcal{B}_{\text{val}}$, keeping all continuations of a branch point in the same set since they share a prefix. For each layer ℓ, we compute $\mathbf{r}_\ell$ with Equation (1) on $\mathcal{B}_{\text{fit}}$, project every held-out continuation onto it, $\langle \mathbf{s}_\ell(c), \mathbf{r}_\ell \rangle$, and measure the ROC AUC for separating terminating from non-terminating continuations. We select $\ell^\star = \arg\max_\ell \mathrm{AUC}(\ell)$ and build the steering vector $\mathbf{v}$ from $\mathbf{r}_{\ell^\star}$ as described in Section 4.3. Algorithm 1 summarizes the full procedure.

		MATH-500 L4–5					AIME-25					OpenR1-Math				
Model	**Ckpt**	**Prob.**	**BP**	**POS**	**NEG**	**Len**	**Prob.**	**BP**	**POS**	**NEG**	**Len**	**Prob.**	**BP**	**POS**	**NEG**	**Len**
Deepseek-R1-Distill -Qwen-1.5B	obc_40	91	91	2780	132	1742	8	14	439	9	1515	112	120	3433	407	2585
	obc_50			2545	367	1575			410	38	1347			2948	892	2484
	obc_60			953	1959	2695			260	188	2076			1034	2806	3392
	obc_70			0	2912	—			0	448	—			0	960	—
Deepseek-R1-Distill -Qwen-7B	obc_40	80	80	2554	6	1577	8	13	415	1	1822	72	103	3229	67	2729
	obc_50			2504	56	1630			407	9	1870			3012	284	2697
	obc_60			2172	388	2353			355	61	2350			2300	996	3182
	obc_70			893	1667	3203			161	255	2391			226	598	3360
Deepseek-R1-Distill -Qwen-14B	obc_40	72	72	2303	1	1431	11	17	544	0	1368	63	100	6382	18	2060
	obc_50			2302	2	1558			544	0	1566			6293	107	2242
	obc_60			2127	177	1966			532	12	2060			5253	1147	2775
	obc_70			1050	1254	3359			340	204	3844			333	467	3667

Table 10: Branch points and continuation outcomes across pruning ratios. **Prob.** denotes problems with at least one branch point; **BP**, the number of branch points; **POS/NEG**, self-terminating/non-terminating continuations; and **Len**, median positive-continuation length. Despite sharing prefixes from dense traces sufficient for correct answers, higher-pruned checkpoints produce more negative continuations and longer median continuation lengths.

Model	**Ckpt**	**N**	**Branch pts**	**POS**	**NEG**	**Len**
Deepseek-Distill-Llama-8B	obc_50	8	103	810	14	1034
	obc_60	8	64	372	140	1743
OpenThinker3-1.5B	obc_50	8	177	1289	127	3398
OpenThinker3-7B	obc_50	8	179	1381	51	3398
	obc_60	24	179	4132	164	3394

Table 11: Continuation outcomes for DeepSeek-R1-Distill-Llama-8B and OpenThinker3 from their steering calibration sets on OpenR1-Math. **Branch pts**, **POS**, **NEG**, and **Len** denote the number of branch points, positive continuations, negative continuations, and median continuation length, respectively.

		DeepSeek-R1-Distill-Qwen			**DeepSeek-R1-Distill-Llama**	**OpenThinker3**	
		1.5B	7B	14B	8B	1.5B	7B
TGAS	Layer $\ell^\star$	18	20	29	9	20	26
	Dense	1	1	1	2	0.5	1.5
	10%	1.5	1.25	1.25	2	1.5	1.5
	50%	1.5	0.5	0.5	1.25	1	1.25
	60%	1.75	1.25	0.75	1.5	1	1.5
	70%	2	1.5	1	2	0.5	1.25
SEAL	Layer ℓ^*	20	20	34	23	20	20
	Dense	1	1	0.5	1	1	0.5
	10%	1	1	0.5	0.5	1	0.25
	50%	1	0.5	0.5	0.5	1	1
	60%	0.5	1	0.25	1	1	0.5
	70%	1	0.5	1	0.5	1.25	0.5

Table 12: Hyperparameters for TGAS and SEAL. ℓ^* denotes the decoder layers. TGAS strength α is selected from $0.5 \rightarrow 2.0$ per checkpoint, and SEAL strength from $0.25 \rightarrow 1.5$ based on performance on the held-out OpenR1-Math set.

C.4.6 Hyperparameter Tuning

TGAS and SEAL both use activation steering and are calibrated on the OpenR1-Math set. TGAS has two hyperparameters: the steering layer $\ell^\star$ and strength α. We select $\ell^\star$ using calibration rollouts on the medium pruning checkpoints with TGAS and on dense model with SEAL and reuse it across all pruned checkpoints, while α is selected per checkpoint based on accuracy on the held-out OpenR1-Math split. We also compare with directions created from medium pruning checkpoints of SEAL at Appendix F.8. The resulting hyperparameters for all models and sparsity levels are reported in Table 12.

Algorithm 1 Training-free Post-pruning via Termination-Guided Activation Steering (TGAS)

Input: Calibration checkpoint M_c (intermediate sparsity) and target checkpoint M_p of the same model; dense correct traces $\mathcal{D}_{\text{cov}}$; number of continuations N; layers $\{1, \ldots, L\}$;
candidate strengths $\mathcal{A}$; held-out OpenR1-Math split $\mathcal{D}_{\text{val}}$.
Output: Steered generation of M_p with steering vector $\mathbf{v}$ at layer $\ell^\star$.

Step 1: Construct branch-point continuation sets {once per model, with M_c}
for each trace in $\mathcal{D}_{\text{cov}}$ **do**
  Identify branch point b: first paragraph start after *Conclusive* s.t. $b < t_{\text{end}}$
  Teacher-force prefix $\mathbf{x}_{<b}$ through M_c; sample N continuations
  Assign continuations to $\mathcal{C}_b^+$ (terminating) or $\mathcal{C}_b^-$ (non-terminating)
end for
Step 2: Compute difference-in-means direction
for $\ell = 1, \ldots, L$ **do**
  for each branch point $b \in \mathcal{B}$ **do**
    for each continuation $c \in \mathcal{C}_b^+ \cup \mathcal{C}_b^-$ **do**
      $\mathbf{s}_\ell(c) \leftarrow \big(\sum_{t=b}^{t_{\text{end}}-1} \mathbf{h}_{\ell,t}(c)\big)/(t_{\text{end}} - b)$ {up to the continuation's end if it is shorter}
    end for
    $\boldsymbol{\mu}_{\ell,b}^+ \leftarrow \big(\sum_{c \in \mathcal{C}_b^+} \mathbf{s}_\ell(c)\big)/|\mathcal{C}_b^+|, \quad \boldsymbol{\mu}_{\ell,b}^- \leftarrow \big(\sum_{c \in \mathcal{C}_b^-} \mathbf{s}_\ell(c)\big)/|\mathcal{C}_b^-|$
  end for
  $\mathbf{r}_\ell \leftarrow \frac{1}{|\mathcal{B}|} \sum_{b \in \mathcal{B}} \big(\boldsymbol{\mu}_{\ell,b}^+ - \boldsymbol{\mu}_{\ell,b}^-\big)$
end for
Step 3: Build the steering vector
$\ell^\star \leftarrow \arg\max_\ell \ \mathrm{AUC}(\mathbf{r}_\ell)$ on held-out branch points {ROC AUC of the projection}
$\bar{n}_{\ell^\star} \leftarrow$ mean ℓ_2 norm of $\mathbf{h}_{\ell^\star,t}$ over calibration trajectories
$\mathbf{v} \leftarrow \bar{n}_{\ell^\star}\, \mathbf{r}_{\ell^\star}/\|\mathbf{r}_{\ell^\star}\|_2$ {reused for all sparsity levels of the model}
$\alpha \leftarrow \arg\max_{\alpha' \in \mathcal{A}} \mathrm{Acc}(M_p, \alpha'\mathbf{v})$ on $\mathcal{D}_{\text{val}}$ {per checkpoint}
Step 4: Apply steering at inference
for each token position t during autoregressive generation **do**
  $\tilde{\mathbf{h}}_{\ell^\star,t} \leftarrow \mathbf{h}_{\ell^\star,t} + \alpha\, \mathbf{v}$
end for
return steered outputs of M_p

D Complete Observation Results

D.1 Answer Retention and Termination in Pruned Models on Dense-Model-Correct Samples

To separate loss of reasoning capability from loss of termination, we restrict evaluation to generations that the dense model answers correctly and track what happens to them after pruning. Table 13 measures how many of these generations the pruned model still solves. Table 14 then asks why the remaining generations fail, classifying each as correct, non-termination, or terminated-wrong, and reports accuracy among terminated generations.

As sparsity increases, pruned models retain fewer problems solved by their dense counterparts, detailed in Table 13, primarily due to increased non-termination rather than incorrect terminated generations, detailed in Table 14. For DeepSeek-R1-Distill-Qwen-7B on AIME-25, non-termination rises from 9% at 10% sparsity to 73% at 70%, while terminated-wrong remains between 14% and 28%. This pattern generalizes across models and benchmarks with non-termination exceeding terminated-wrong rates in 10 of 12 pairs at 60% sparsity. Moreover, conditional on termination, larger models retain relatively high accuracy. For example, DeepSeek-R1-Distill-Qwen-14B achieves 73.8% and 83.3%, while OpenThinker3-7B achieves 74.4% and 63.9% on AIME-25 and LiveCodeBench, respectively. These results suggest that pruning disproportionately impairs termination behavior, while the ability to reason correctly remains comparatively robust when generation terminates.

D.2 Comprehensive Evaluation of Larger Budget Generation

Increasing the generation limit from 16,384 to 65,536 tokens provides limited benefit for non-terminating traces, detailed in Table 15. While the extended budget substantially reduces non-termination for dense models, $20/24$ of the 7B and $15/15$ of the 14B non-terminating traces terminate under the larger budget, this yields only 0 and 2 additional correct answers, respectively. Recovery becomes increasingly rare under stronger pruning: for 60% pruned by OBC, only $10/68$ of 7B model and $4/73$ of 14B model traces are recovered, while 70% pruned recovers none. This suggests that insufficient generation length is not the primary cause of failure; rather, pruning induces a persistent tendency to continue generating instead of committing to a final answer.

D.3 Comprehensive Evaluation of the Stop-Token Mechanism

We evaluate the pruned checkpoints using forward passes with teacher forcing along the dense model's trajectories. We focus on the $\mathcal{D}_{\text{cov}}$ set, where the dense models terminate and answer correctly. Each trace is replayed under its original generation prompt. We focus on two position sets: **STOP**, the position where the dense model emits `</think>`, and **OTHER**, positions more than 50 tokens before STOP. At each position, we measure the probability and rank assigned to `</think>`. We define Δ_{close} as the log-probability of `</think>` at STOP minus its mean log-probability over OTHER positions, and Δ_{stop} as the log-probability of `</think>` at STOP minus the mean log-probability of the token the dense model actually generated at OTHER positions. The full results are reported in Table 16.

Across models and tasks, the stopping signal remains highly localized at the dense model's STOP position. For the 7B and 14B models, the median rank of `</think>` is 0 at 10%–60% pruning ratios, meaning that `</think>` is consistently the top prediction at the correct stopping position. Even at 70% sparsity, where this signal weakens, `</think>` remains near the top: its median rank is 5.5 and 2 on AIME-25 for the 7B and 14B models, respectively. The same pattern holds on MATH-500, with median ranks of 7 and 3. At OTHER positions, in contrast, `</think>` receives almost no probability at any checkpoint: its mean log-probability lies between -23.3 and -30.4 for the 7B and 14B models.

The log-probability of `</think>` at STOP is substantially higher than at OTHER positions. For the 7B model, $\log p(\texttt{</think>})$ at STOP rises from -4.83 at 70% pruning ratio to near zero at pruned checkpoints 10%–60%, while remaining between -23 and -28 at OTHER positions. Consequently, Δ_{close} remains at least $+22$ for both 7B and 14B models at 70% sparsity. Meanwhile, Δ_{stop} is positive below 70% sparsity but drops to -2.6 to -4.3 at 70%, indicating that aggres-

		AIME25		LCB		MATH-500 L4–5	
family	level	n kept	kept	n kept	kept	n kept	kept
DeepSeek-R1-Distill-Qwen-1.5B (OBC)	dense	27.0 ± 0.0	100	133.0 ± 2.8	100	197.0 ± 1.4	100
	sparse10	19.5 ± 0.7	72.2 ± 2.6%	107.5 ± 3.5	80.8 ± 0.9%	175.0 ± 2.8	88.8 ± 2.1%
	sparse50	15.0 ± 1.4	55.6 ± 5.2%	51.0 ± 4.2	38.3 ± 2.4%	133.5 ± 6.4	67.8 ± 2.7%
	sparse60	1.5 ± 0.7	5.6 ± 2.6%	0.0 ± 0.0	0.0 ± 0.0%	39.5 ± 9.2	20.0 ± 4.5%
DeepSeek-R1-Distill-Qwen-7B (OBC)	dense	47.0 ± 5.7	100	223.5 ± 2.1	100	223.0 ± 2.8	100
	sparse10	33.0 ± 1.4	70.5 ± 5.5%	198.0 ± 2.8	88.6 ± 2.1%	208.0 ± 4.2	93.3 ± 0.7%
	sparse50	30.0 ± 1.4	64.1 ± 4.7%	145.0 ± 7.1	64.9 ± 2.5%	204.5 ± 0.7	91.7 ± 0.8%
	sparse60	21.0 ± 4.2	44.5 ± 3.7%	55.5 ± 0.7	24.8 ± 0.6%	173.5 ± 0.7	77.8 ± 0.7%
	sparse70	6.0 ± 0.0	12.9 ± 1.5%	0.0 ± 0.0	0.0 ± 0.0%	57.0 ± 1.4	25.6 ± 1.0%
DeepSeek-R1-Distill-Qwen-14B (OBC)	dense	64.0 ± 1.4	100	278.0 ± 8.6	100	231.5 ± 0.7	100
	sparse10	47.5 ± 0.7	74.2 ± 0.5%	254.0 ± 4.7	91.4 ± 1.7%	220.0 ± 4.2	95.0 ± 1.5%
	sparse50	37.5 ± 0.7	58.6 ± 0.2%	225.0 ± 6.5	80.9 ± 2.4%	206.5 ± 2.1	89.2 ± 0.6%
	sparse60	29.5 ± 0.7	46.1 ± 0.1%	140.0 ± 8.3	50.4 ± 3.0%	172.0 ± 7.1	74.3 ± 2.8%
	sparse70	3.0 ± 1.4	4.7 ± 2.1%	0.0 ± 0.0	0.0 ± 0.0%	40.5 ± 2.1	17.5 ± 1.0%
DeepSeek-R1-Distill-Llama-8B (OBC)	dense	10.2 ± 1.7	100	55.8 ± 2.4	100	193.0 ± 7.1	100
	sparse10	7.8 ± 1.7	76.5 ± 17.3%	49.8 ± 2.9	89.3 ± 5.9%	178.0 ± 3.7	92.2 ± 1.9%
	sparse40	7.2 ± 1.0	73.2 ± 20.6%	47.0 ± 2.2	84.3 ± 0.8%	155.0 ± 5.5	80.3 ± 2.9%
	sparse50	5.5 ± 1.7	54.2 ± 16.2%	38.8 ± 1.0	69.6 ± 2.1%	119.0 ± 6.8	61.7 ± 3.5%
	sparse60	1.0 ± 0.8	9.6 ± 7.1%	10.2 ± 1.5	18.4 ± 3.1%	31.0 ± 5.1	16.1 ± 2.6%
	sparse70	0.0 ± 0.0	0.0 ± 0.0%	0.0 ± 0.0	0.0 ± 0.0%	0.0 ± 0.0	0.0 ± 0.0%
OpenThinker3-1.5B (OBC)	dense	9.8 ± 0.5	100	41.0 ± 2.7	100	199.0 ± 6.9	100
	sparse10	7.8 ± 1.0	79.4 ± 8.2%	35.2 ± 2.2	86.0 ± 1.8%	172.0 ± 4.8	86.4 ± 2.4%
	sparse40	6.2 ± 1.0	64.4 ± 12.1%	25.5 ± 2.5	62.2 ± 3.7%	139.0 ± 6.5	69.8 ± 3.3%
	sparse50	5.2 ± 0.5	53.9 ± 4.8%	8.5 ± 0.6	20.8 ± 1.2%	69.0 ± 6.7	34.7 ± 3.4%
	sparse60	0.2 ± 0.5	2.8 ± 5.6%	0.0 ± 0.0	0.0 ± 0.0%	7.0 ± 2.6	3.5 ± 1.3%
OpenThinker3-7B (OBC)	dense	13.2 ± 1.5	100	60.2 ± 1.9	100	198.0 ± 7.0	100
	sparse10	11.5 ± 1.3	87.1 ± 9.2%	53.8 ± 1.7	89.3 ± 4.8%	177.0 ± 4.3	89.4 ± 2.2%
	sparse40	11.0 ± 1.8	82.7 ± 5.3%	53.2 ± 1.0	88.4 ± 2.4%	163.0 ± 5.4	82.3 ± 2.7%
	sparse50	9.8 ± 0.5	74.4 ± 10.4%	46.0 ± 1.6	76.3 ± 1.4%	157.0 ± 5.7	79.3 ± 2.9%
	sparse60	7.0 ± 1.6	52.5 ± 8.4%	28.8 ± 1.7	47.7 ± 1.8%	126.0 ± 6.8	63.6 ± 3.4%
	sparse70	0.8 ± 0.5	6.0 ± 4.0%	0.0 ± 0.0	0.0 ± 0.0%	1.0 ± 1.0	0.5 ± 0.5%

Table 13: Dense-correct generations retained by the pruned model. n *kept*: correct for both models; *kept*: fraction of dense-correct generations also solved by the pruned model. Values are mean ± std.

sive pruning weakens the stopping signal relative to ordinary next-token prediction. This pattern holds across AIME-25, MATH-500, and both model sizes, suggesting that pruning preserves a localized stopping signal even at high sparsity.

D.4 Comprehensive Evaluation of Transition of Pruned LRMs

We investigate whether pruning alters reasoning trajectories at moments of reconsideration, where a model may either advance its solution or revisit previously explored reasoning. We define a *Decision point* as the start of a paragraph beginning with one of 11 reconsideration cues: *wait*, *but wait*, *hmm*, *alternatively*, *actually*, *hold on*, or *let me check/verify/double-check/reconsider/recheck*. Cues are matched case-insensitively at the beginning of `\n\n`-separated paragraphs, followed Lotfi et al. (2026). The cue set is fixed beforehand using paragraph-initial bigram frequencies from dense-model traces and is not modified based on pruned generations. *Control points* are randomly sampled paragraph starts without these cues. We also assign each paragraph a local trajectory label: CYCLIC if the Jaccard similarity between its set of 8-word n-grams and that of any preceding paragraph in the same trace is at least 0.5, and PROGRESS otherwise.

Teacher-Forced Preferences. We compute Δ_{cyc} (Section 3) at decision and control points to assess whether pruning increases the tendency toward CYCLIC reasoning at reconsideration points and ordinary paragraph starts. As shown in Table 17, Δ_{cyc} becomes progressively less negative as pruning increases across all model families and tasks, indicating a growing preference for cyclic over progressive continuations. This pattern also holds at control points, indicating that pruning broadly shifts continuation preferences rather than selectively affecting reconsideration cues. The effect is strongest at high sparsity, whereas predictive entropy changes more gradually, suggesting a shift in continuation behavior rather than a comparable increase in uncertainty.

Free-Running Continuations. We examine how pruned models continue dense reasoning traces whose prefixes are sufficient to identify the correct answer. For each eligible trace, we teacher-force the prefix up to the branch point

		AIME25				LCB				MATH-500 L4–5			
family	level	(a)	(b)	(c)	kept | term.	(a)	(b)	(c)	kept | term.	(a)	(b)	(c)	kept | term.
DeepSeek-R1-Distill-Qwen-1.5B (OBC)	sparse10	19.5 (72)	2.5 (9)	5.0 (19)	$79.7 \pm 5.2\%$	107.5 (81)	4.5 (3)	21.0 (16)	$83.7 \pm 0.5\%$	175.0 (89)	5.5 (3)	16.5 (8)	$91.4 \pm 0.5\%$
	sparse50	15.0 (56)	6.5 (24)	5.5 (20)	$73.1 \pm 4.4\%$	51.0 (38)	50.5 (38)	31.5 (24)	$61.8 \pm 1.4\%$	133.5 (68)	38.0 (19)	25.5 (13)	$83.9 \pm 3.3\%$
	sparse60	1.5 (6)	22.5 (83)	3.0 (11)	*35.0 ± 21.2%*	0.0 (0)	132.5 (100)	0.5 (0)	*0.0%*	39.5 (20)	142.5 (72)	15.0 (8)	$72.0 \pm 8.5\%$
DeepSeek-R1-Distill-Qwen-7B (OBC)	sparse10	33.0 (71)	4.0 (9)	10.0 (21)	$77.5 \pm 9.5\%$	198.0 (89)	4.5 (2)	21.0 (9)	$90.4 \pm 1.3\%$	208.0 (93)	3.0 (1)	12.0 (5)	$94.5 \pm 1.3\%$
	sparse50	30.0 (64)	8.5 (18)	8.5 (18)	$78.1 \pm 3.5\%$	145.0 (65)	25.5 (11)	53.0 (24)	$73.2 \pm 1.0\%$	204.5 (92)	8.5 (4)	10.0 (4)	$95.3 \pm 0.6\%$
	sparse60	21.0 (44)	13.0 (28)	13.0 (28)	$61.6 \pm 2.2\%$	55.5 (25)	108.0 (48)	60.0 (27)	$48.1 \pm 0.3\%$	173.5 (78)	37.5 (17)	12.0 (5)	$93.5 \pm 0.7\%$
	sparse70	6.0 (13)	34.5 (73)	6.5 (14)	*48.1 ± 2.7%*	0.0 (0)	225.0 (100)	0.0 (0)	–	57.0 (26)	150.5 (67)	15.5 (7)	$78.6 \pm 2.7\%$
DeepSeek-R1-Distill-Qwen-14B (OBC)	sparse10	47.5 (74)	5.5 (9)	11.0 (17)	$81.2 \pm 1.7\%$	254.0 (91)	5.0 (2)	19.0 (7)	$93.0 \pm 1.5\%$	220.0 (95)	1.5 (1)	10.0 (4)	$95.6 \pm 0.7\%$
	sparse50	37.5 (59)	11.5 (18)	15.0 (23)	$71.4 \pm 0.4\%$	225.0 (81)	22.0 (8)	31.0 (11)	$87.9 \pm 2.0\%$	206.5 (89)	5.5 (2)	19.5 (8)	$91.4 \pm 0.8\%$
	sparse60	29.5 (46)	24.0 (37)	10.5 (16)	$73.8 \pm 4.4\%$	140.0 (50)	110.0 (40)	28.0 (10)	$83.3 \pm 2.9\%$	172.0 (74)	38.0 (16)	21.5 (9)	$88.9 \pm 1.4\%$
	sparse70	3.0 (5)	53.5 (84)	7.5 (12)	*27.9 ± 4.1%*	0.0 (0)	278.0 (100)	0.0 (0)	–	40.5 (17)	172.5 (75)	18.5 (8)	$68.6 \pm 2.0\%$
DeepSeek-R1-Distill-Llama-8B (OBC)	sparse10	7.8 (77)	0.8 (6)	1.8 (17)	$81.4 \pm 12.6\%$	49.8 (89)	1.2 (2)	4.8 (8)	$91.3 \pm 5.0\%$	178.0 (92)	7.0 (4)	8.0 (4)	$95.7 \pm 1.5\%$
	sparse40	7.2 (73)	0.5 (4)	2.5 (22)	$76.1 \pm 19.0\%$	47.0 (84)	1.8 (3)	7.0 (13)	$87.1 \pm 1.2\%$	155.0 (80)	28.0 (15)	10.0 (5)	$93.9 \pm 1.9\%$
	sparse50	5.5 (54)	0.8 (6)	4.0 (39)	$57.5 \pm 15.0\%$	38.8 (70)	7.2 (13)	9.8 (17)	$80.2 \pm 5.4\%$	119.0 (62)	63.0 (33)	11.0 (6)	$91.5 \pm 2.4\%$
	sparse60	1.0 (10)	6.2 (60)	3.0 (30)	*28.8 ± 21.7%*	10.2 (18)	33.0 (59)	12.5 (22)	$44.9 \pm 4.7\%$	31.0 (16)	150.0 (78)	12.0 (6)	$72.1 \pm 6.8\%$
	sparse70	0.0 (0)	10.2 (100)	0.0 (0)	–	0.0 (0)	55.8 (100)	0.0 (0)	–	0.0 (0)	193.0 (100)	0.0 (0)	–
OpenThinker3-1.5B (OBC)	sparse10	7.8 (79)	1.0 (10)	1.0 (10)	$88.5 \pm 9.1\%$	35.2 (86)	2.5 (6)	3.2 (8)	$91.6 \pm 3.8\%$	172.0 (86)	21.0 (11)	6.0 (3)	$96.6 \pm 1.4\%$
	sparse40	6.2 (64)	3.0 (31)	0.5 (5)	*92.9 ± 14.3%*	25.5 (62)	7.8 (19)	7.8 (19)	$76.6 \pm 5.8\%$	139.0 (70)	51.0 (26)	9.0 (5)	$93.9 \pm 2.0\%$
	sparse50	5.2 (54)	3.2 (34)	1.2 (12)	*82.1 ± 13.7%*	8.5 (21)	21.8 (53)	10.8 (26)	$44.4 \pm 2.6\%$	69.0 (35)	122.0 (61)	8.0 (4)	$89.6 \pm 3.5\%$
	sparse60	0.2 (3)	8.5 (87)	1.0 (10)	*12.5 ± 25.0%*	0.0 (0)	40.5 (99)	0.5 (1)	*0.0 ± 0.0%*	7.0 (4)	189.0 (95)	3.0 (2)	*70.0 ± 14.5%*
OpenThinker3-7B (OBC)	sparse10	11.5 (87)	1.0 (7)	0.8 (6)	$93.9 \pm 4.2\%$	53.8 (89)	3.0 (5)	3.5 (6)	$93.9 \pm 3.5\%$	177.0 (89)	3.0 (2)	18.0 (9)	$90.8 \pm 2.1\%$
	sparse40	11.0 (83)	2.0 (16)	0.2 (2)	$98.2 \pm 3.6\%$	53.2 (88)	3.8 (6)	3.2 (5)	$94.3 \pm 2.4\%$	163.0 (82)	11.0 (6)	24.0 (12)	$87.2 \pm 2.4\%$
	sparse50	9.8 (74)	2.0 (15)	1.5 (10)	$88.0 \pm 13.9\%$	46.0 (76)	8.8 (15)	5.5 (9)	$89.4 \pm 3.2\%$	157.0 (79)	26.0 (13)	15.0 (8)	$91.3 \pm 2.2\%$
	sparse60	7.0 (52)	3.8 (28)	2.5 (20)	$74.4 \pm 13.2\%$	28.8 (48)	15.2 (25)	16.2 (27)	$63.9 \pm 2.1\%$	126.0 (64)	59.0 (30)	13.0 (7)	$90.6 \pm 2.5\%$
	sparse70	0.8 (6)	10.5 (79)	2.0 (15)	*43.8 ± 42.7%*	0.0 (0)	59.8 (99)	0.5 (1)	*0.0 ± 0.0%*	1.0 (1)	196.0 (99)	1.0 (1)	*50.0 ± 35.4%*

Table 14: Breakdown of dense-correct generations under pruning. Each generation is (a) *correct*, (b) *reasoning divergence* (no </think>), or (c) *terminated-wrong* (</think> emitted but the final answer is incorrect). Values show mean count per run with the corresponding percentage in parentheses. *Kept | term.* denotes accuracy among terminated generations. Italics indicate fewer than 30 terminated dense-correct items.

through the pruned checkpoint and sample N continuations c.

Table 18 shows consistent shifts toward cyclic reasoning across all models and benchmarks. Loop entry, $P(\text{cyc} \mid \text{prog})$, increases monotonically with sparsity, reaching 0.097-0.105 on MATH-500 and 0.081-0.085 on AIME-25 for DS-Qwen models at 70% sparsity, roughly 3–6× their dense rates. Loop persistence, $P(\text{cyc} \mid \text{cyc})$, also rises, though it peaks at 60% in three of four DS-Qwen blocks. OT3-7B exhibits a weaker but similar increase, particularly in persistence. These trends accompany a 3–5× increase in paragraph pairs from 40% to 70% sparsity, suggesting that pruning promotes both entry into and persistence within cyclic reasoning across model families and benchmarks.

D.5 Comprehensive Commitment Directions Analysis

Section 4.2 shows that termination directions constructed from different pruned checkpoints of the same model have high cosine similarity. However, since these checkpoints are pruned from the same dense model, some similarity is expected regardless of what the directions encode (Ding et al., 2021). We therefore use hypothesis testing with three increasingly strong baselines and a transfer test to assess whether the observed cosine similarity is significant and specifically associated with termination.

Setup. We consider all pairs of sparsity levels (A, B) within the same model for which both $\mathcal{C}^+$ and $\mathcal{C}^-$ contain sufficient continuations, yielding 15 checkpoint pairs across four models on MATH-500 and OpenR1-Math. All checkpoints share the same branch points, with 32 continuations sampled per point (64 for DS-Qwen-14B on OpenR1-Math), using temperature 0.6 and a 16,384-token budget. To make directions comparable, we use the same number N of terminating and non-terminating continuations at every sparsity level of a model, and draw them from the same set of problems at every level. Thus, N is determined by the scarcest class and ranges from 52 to 200. Let $\mathbf{r}_\ell^A$ denote the direction at layer ℓ for checkpoint A. We measure similarity using the layer-averaged cosine:

$$\text{sim}(A, B) = \frac{1}{L} \sum_{\ell=1}^{L} \cos\big(\mathbf{r}_\ell^A, \mathbf{r}_\ell^B\big).$$

All baselines reuse the same continuation representations $\mathbf{s}_\ell(c)$ and layers, varying only the labels or held-out split. We next compare the cosine similarities against baseline directions:

Random directions. Two independent random unit vectors in d dimensions have an expected ab-

Model	Level	Token Limit	Non-termination Rate	Accuracy	Accuracy | Terminated	Recovered Traces	Recovered Correct
7B	dense	16,384	20.0%	42.5%	53.1%	—	—
7B	dense	65,536	3.3%	42.5%	44.0%	20 / 24 (83.3%)	0
7B	obc_10	16,384	20.0%	31.7%	39.6%	—	—
7B	obc_10	65,536	2.5%	31.7%	32.5%	21 / 24 (87.5%)	0
7B	obc_50	16,384	42.5%	30.0%	52.2%	—	—
7B	obc_50	65,536	12.5%	30.8%	35.2%	36 / 51 (70.6%)	2
7B	obc_60	16,384	56.7%	22.5%	51.9%	—	—
7B	obc_60	65,536	48.3%	22.5%	43.5%	10 / 68 (14.7%)	0
7B	obc_70	16,384	80.8%	5.0%	26.1%	—	—
7B	obc_70	65,536	80.8%	5.0%	26.1%	0 / 97 (0.0%)	0
14B	dense	16,384	12.5%	54.2%	61.9%	—	—
14B	dense	65,536	0.0%	55.8%	55.8%	15 / 15 (100.0%)	2
14B	obc_10	16,384	15.8%	47.5%	56.4%	—	—
14B	obc_10	65,536	0.0%	49.2%	49.2%	19 / 19 (100.0%)	2
14B	obc_50	16,384	36.7%	35.8%	56.6%	—	—
14B	obc_50	65,536	10.8%	38.3%	43.0%	31 / 44 (70.5%)	3
14B	obc_60	16,384	60.8%	26.7%	68.1%	—	—
14B	obc_60	65,536	57.5%	27.5%	64.7%	4 / 73 (5.5%)	1
14B	obc_70	16,384	89.2%	3.3%	30.8%	—	—
14B	obc_70	65,536	89.2%	3.3%	30.8%	0 / 107 (0.0%)	0

Table 15: Effect of increasing the generation budget from 16,384 to 65,536 tokens on the DeepSeek-R1-Distill-Qwen 7B and 14B models on AIME-25. A 65,536 row continues each generation non-termination at 16,384 from its own tokens; *recovered* = non-termination at 16,384 and terminated by 65,536.

		AIME-25						MATH-500 L4–5					
Model	Ckpt	p(</think>)	Med. rank	logp STOP	logp </think> OTHER	Δ_{close}	Δ_{stop}	p(</think>)	Med. rank	logp STOP	logp </think> OTHER	Δ_{close}	Δ_{stop}
Deepseek-R1-Distill -Qwen-1.5B	obc_10	*0.9999*	0	−0.000	−22.71	+22.71	+0.377	0.9999	0	−0.000	−23.11	+23.11	+0.374
	obc_50	*0.9804*	0	−0.020	−27.28	+27.26	+0.400	0.9918	0	−0.008	−26.51	+26.50	+0.428
	obc_60	*0.7273*	0	−0.379	−26.56	+26.18	+0.134	0.6301	0	−0.771	−26.06	+25.29	−0.226
	obc_70	*0.0351*	44	−9.192	−24.19	+14.99	−8.431	0.0154	157.5	−11.076	−23.84	+12.76	−10.229
Deepseek-R1-Distill -Qwen-7B	obc_10	*1.0000*	0	−0.000	−23.30	+23.30	+0.361	0.9999	0	−0.000	−24.15	+24.15	+0.349
	obc_50	*0.9978*	0	−0.002	−25.23	+25.23	+0.374	0.9963	0	−0.004	−25.50	+25.50	+0.375
	obc_60	*0.9570*	0	−0.045	−27.75	+27.71	+0.378	0.9240	0	−0.097	−27.03	+26.93	+0.345
	obc_70	*0.0360*	5.5	−4.826	−27.43	+22.61	−4.274	0.0400	7	−4.747	−26.75	+22.00	−4.133
Deepseek-R1-Distill -Qwen-14B	obc_10	*1.0000*	0	−0.000	−24.68	+24.68	+0.353	1.0000	0	−0.000	−25.47	+25.47	+0.340
	obc_50	*0.9945*	0	−0.006	−28.45	+28.44	+0.365	0.9976	0	−0.002	−27.61	+27.61	+0.376
	obc_60	*0.8539*	0	−0.223	−30.35	+30.13	+0.212	0.9069	0	−0.130	−28.96	+28.83	+0.331
	obc_70	*0.0725*	2	−3.748	−28.06	+24.31	−3.158	0.1304	3	−3.298	−26.95	+23.65	−2.642
OpenThinker3 -7B	obc_10	*1.0000*	0	−0.000	−31.12	+31.12	+0.319	0.9972	0	−0.004	−31.89	+31.89	+0.291
	obc_50	*0.9994*	0	−0.001	−29.71	+29.71	+0.393	0.9924	0	−0.021	−30.15	+30.13	+0.368
	obc_60	*0.9978*	0	−0.002	−28.81	+28.80	+0.467	0.9826	0	−0.039	−29.12	+29.08	+0.442
	obc_70	*0.8391*	0	−0.279	−26.96	+26.68	+0.371	0.5784	0	−0.912	−27.05	+26.14	−0.210

Table 16: Teacher-forced evaluation of the stopping signal at the dense model's </think> position. STOP denotes the position whose next token is the dense </think>, while OTHER includes all trace positions more than 50 tokens before STOP. Logp STOP is the mean log-probability of </think> at STOP. Logp </think> OTHER is the mean log-probability of </think> over OTHER positions, and Δ_{close} = logp STOP − logp </think> OTHER. Δ_{stop} = logp STOP − the mean log-probability of the token the dense model actually generated at OTHER positions. Rank 0 indicates that </think> is the model's top prediction at STOP. All quantities are computed by teacher-forcing dense-model traces through each checkpoint.

solute cosine of $\sqrt{2/(\pi d)}$, at most 0.020 for our models (Table 19). All observed similarities substantially exceed this random baseline.

Label-shuffled directions. We first test whether any difference of means over these continuations aligns across checkpoints, regardless of the labels. For each pair, we randomly permute the terminating and non-terminating labels of checkpoint A while keeping the class sizes, recompute its direction at every layer, and measure its similarity with the true direction of B, repeated 200 times. As shown in Table 20, shuffled directions already reach 0.32–0.43 on average and 0.68–0.82 at the 95th percentile, far above random directions. Nevertheless, the real similarity exceeds the shuffled baseline in all 15 pairs ($p \leq 0.030$). The similarity therefore depends on the termination labels.

A shared contrast unrelated to termination. The label-shuffled baseline does not rule out generic alignment between checkpoints pruned from the same model. As a control, we use branch-point identity, which all checkpoints encode since they share the same branch points, but which is unrelated to termination.

Concretely, we draw a random half $\mathcal{H} \subset \mathcal{B}$ of the branch points and relabel every continuation of A and B by whether its branch point lies in $\mathcal{H}$, ignoring whether it terminates. Using the same $\mathcal{H}$

		AIME-25					LCB					MATH-500 L4–5				
Model	**Ckpt**	n_{dec}	n_{ctrl}	**Dec.**	**Ctrl.**	**Ent.**	n_{dec}	n_{ctrl}	**Dec.**	**Ctrl.**	**Ent.**	n_{dec}	n_{ctrl}	**Dec.**	**Ctrl.**	**Ent.**
DS-Qwen 1.5B	dense	796	200	−13.0	−13.2	1.60	720	200	−16.7	−15.5	1.46	684	200	−14.9	−21.9	1.55
	obc_10	794	200	−13.4	−13.4	1.62	734	200	−16.0	−12.5	1.44	690	200	−15.3	−23.3	1.57
	obc_50	786	194	−8.4	−11.7	1.32	788	192	−8.2	−5.3	1.48	686	200	−10.2	−15.3	1.35
	obc_60	762	182	−3.8	−4.5	1.33	756	172	−2.7	−3.2	1.33	762	198	−5.6	−6.8	1.32
	obc_70	524	154	1.6	−1.2	1.48	556	158	2.2	−0.9	1.60	506	160	1.0	−0.8	1.54
DS-Qwen 7B	dense	794	200	−15.3	−16.6	1.55	692	200	−19.5	−17.0	1.33	624	200	−18.2	−24.4	1.49
	obc_10	792	200	−14.2	−15.6	1.59	708	200	−18.8	−14.7	1.33	638	200	−18.2	−24.7	1.48
	obc_50	788	200	−11.9	−15.0	1.51	742	200	−14.8	−11.8	1.37	680	200	−15.3	−21.0	1.52
	obc_60	796	200	−9.6	−10.1	1.41	794	200	−10.4	−6.5	1.45	710	200	−12.0	−14.9	1.42
	obc_70	782	188	−5.0	−4.6	1.44	742	182	−4.1	−1.5	1.33	782	194	−6.7	−4.6	1.45
DS-Qwen 14B	dense	790	200	−13.8	−16.0	1.54	680	200	−17.1	−18.1	1.33	622	200	−16.1	−23.7	1.41
	obc_10	788	200	−13.2	−15.8	1.49	670	200	−17.5	−17.6	1.26	630	200	−16.0	−23.2	1.47
	obc_50	768	198	−12.4	−15.9	1.40	680	200	−13.5	−13.5	1.30	650	200	−14.7	−21.6	1.34
	obc_60	782	196	−10.7	−10.5	1.23	710	199	−9.2	−8.7	1.33	690	196	−12.7	−12.5	1.28
	obc_70	762	180	−4.7	−4.3	1.24	720	174	−3.2	−2.7	1.39	782	196	−5.9	−5.1	1.32
OT3-7B	dense	800	200	−23.1	−23.8	1.03	704	198	−28.7	−21.6	1.17	710	200	−29.3	−37.0	0.88
	obc_10	796	200	−22.6	−24.3	1.09	716	200	−26.4	−22.7	1.17	724	200	−29.5	−30.6	0.88
	obc_40	800	200	−19.6	−18.4	1.14	736	200	−21.6	−20.1	1.30	730	200	−25.5	−27.2	0.96
	obc_50	800	200	−15.8	−17.4	1.22	772	200	−21.9	−15.3	1.33	748	200	−20.1	−25.8	0.99
	obc_60	796	200	−13.2	−12.0	1.33	794	200	−13.1	−10.5	1.51	778	196	−14.6	−17.0	1.11
	obc_70	790	194	−5.6	−5.6	1.60	696	192	−1.4	−3.6	1.81	538	120	−6.7	−12.3	1.47

Table 17: Log-probability ratio between cyclic and progressive continuations on AIME-25, LCB and MATH-500. $\Delta_{\mathrm{cyc}} = \log P(\mathrm{CYCLIC}) - \log P(\mathrm{PROGRESS})$, scored over the next 32 tokens. **Dec.** is measured at decision points, **Ctrl.** at randomly sampled paragraph in the same traces; **Ent.** is the predictive entropy at the decision points. n_{dec} is the number of decision points and n_{ctrl} the number of control points. More negative means the model prefers the progressive continuation more strongly.

		MATH-500 L4–5			AIME-25		
Model	**Ckpt**	P(**cyc** | **cyc**)	P(**cyc** | **prog**)	**Count**	P(**cyc** | **cyc**)	P(**cyc** | **prog**)	**Count**
DeepSeek-R1-Distill-Qwen 7B	dense	0.157	0.025	4094	0.029	0.026	1296
	obc_40	0.206	0.027	155830	0.230	0.031	38947
	obc_50	0.257	0.047	207714	0.283	0.043	51479
	obc_60	0.325	0.086	434020	0.327	0.064	85675
	obc_70	0.307	0.105	596073	0.307	0.085	113353
DeepSeek-R1-Distill-Qwen 14B	dense	0.062	0.017	2714	0.000	0.029	770
	obc_40	0.142	0.028	122288	0.164	0.025	30508
	obc_50	0.212	0.042	180908	0.203	0.033	47349
	obc_60	0.326	0.076	283318	0.271	0.059	68347
	obc_70	0.307	0.097	514252	0.324	0.081	136784
OpenThinker3-7B	dense	0.137	0.026	6227	0.122	0.015	3057
	obc_40	0.138	0.018	306167	0.142	0.018	105101
	obc_50	0.160	0.025	369215	0.146	0.022	127031
	obc_60	0.160	0.032	579485	0.151	0.029	192809
	obc_70	0.210	0.060	1250119	0.167	0.049	479539

Table 18: Transition dynamics between cyclic and progressive states on MATH-500 Levels 4–5 and AIME-25, measured on continuations from a *shared* dense branch point. $P(\mathrm{cyc} \mid \mathrm{cyc})$ and $P(\mathrm{cyc} \mid \mathrm{prog})$ measure loop persistence and loop entry over adjacent paragraph pairs; **Count** is the number of such pairs.

for both checkpoints, we compute the difference-of-means direction of this contrast at every layer for each checkpoint and measure $\mathrm{sim}(A, B)$ between them. We repeat this for 100 random halves and report the average. If shared weights alone explained the alignment, this branch-point contrast would align as strongly as the termination direction. Instead, its similarity is only 0.35–0.81, which is 0.143–0.52 lower than that of the termination direction in all 15 pairs. Thus, the alignment reflects the termination contrast rather than shared weights alone.

Transfer across sparsity levels. Finally, we test whether the shared direction is functional. We split the branch points into five folds, keeping all continuations of a branch point in the same fold since they share a prefix. In each fold, we construct the direction of A from four folds, project the held-out continuations of B onto it, and compute the ROC AUC for separating terminating from non-terminating continuations, averaged over folds and layers. Setting $A = B$ gives each direction's own AUC. As a control, we repeat the transfer with the labels of A shuffled over 200 permutations. All directions achieve high within-checkpoint AUCs (0.914–0.991), while transferred directions on DS-Qwen reach 0.916–0.991. Transfer AUC exceeds the 95th percentile of the shuffled-label control in 29 of 30 pairs. Together with the cosine results, these findings show that the shared direction consistently separates terminating from non-terminating continuations across sparsity levels, as also illustrated by the PCA projections in Figure 10.

In conclusion, our results show that: (i) The co-

Model	d	Random	Observed
DS-Qwen-1.5B	1536	0.020	0.834–0.977
DS-Qwen-7B	3584	0.013	0.870–0.960
OT3-7B	3584	0.013	0.784
DS-Qwen-14B	5120	0.011	0.870–0.881

Table 19: Expected absolute cosine between two random unit vectors, $\sqrt{2/(\pi d)}$, for each model's hidden size d. *Observed*: cosine similarity between termination directions; *Random*: cosine similarity with random directions.

Model	Levels	N	Sim.	Shuffled	p	Split
OpenR1-Math						
DS-Qwen-1.5B	40 / 50	200	0.977	0.40 / 0.77	0.005	0.80
DS-Qwen-1.5B	40 / 60	200	0.878	0.36 / 0.71	0.005	0.68
DS-Qwen-1.5B	50 / 60	200	0.925	0.42 / 0.78	0.005	0.75
DS-Qwen-7B	40 / 50	52	0.951	0.40 / 0.74	0.005	0.79
DS-Qwen-7B	40 / 60	52	0.877	0.35 / 0.69	0.005	0.69
DS-Qwen-7B	50 / 60	52	0.953	0.36 / 0.71	0.005	0.81
DS-Qwen-14B	50 / 60	61	0.881	0.32 / 0.68	0.005	0.69
MATH-500						
DS-Qwen-1.5B	40 / 50	117	0.954	0.42 / 0.79	0.005	0.70
DS-Qwen-1.5B	40 / 60	117	0.834	0.37 / 0.74	0.005	0.58
DS-Qwen-1.5B	50 / 60	117	0.919	0.37 / 0.72	0.005	0.76
DS-Qwen-7B	50 / 60	55	0.960	0.43 / 0.82	0.005	0.59
DS-Qwen-7B	50 / 70	55	0.870	0.41 / 0.76	0.005	0.35
DS-Qwen-7B	60 / 70	55	0.903	0.37 / 0.73	0.005	0.68
DS-Qwen-14B	60 / 70	152	0.870	0.39 / 0.70	0.005	0.58
OT3-7B	60 / 70	126	0.784	0.42 / 0.74	0.030	0.48

Table 20: Cosine similarity between termination directions of two sparsity levels (%) of the same model, against two baselines. N: continuations per class; **Sim.**: layer-averaged cosine similarity; **Shuffled**: mean / 95th percentile of the similarity when the labels of one checkpoint are permuted (200 permutations), with permutation p-value p; **Split**: similarity of a contrast between two random halves of the branch points.

sine similarity of termination directions is consistently higher than that of all ablation baselines; (ii) Directions transfer across sparsity levels, consistently separating terminating from non-terminating continuations; (iii) These findings support learning one termination direction per model and reusing it across sparsity levels.

(a) OpenR1-Math

Model	Levels	Transfer	Own	Shuffled	p
DS-Qwen-1.5B	40 → 50	0.973	0.976	0.839 / 0.950	0.005
DS-Qwen-1.5B	40 → 60	0.974	0.977	0.848 / 0.945	0.005
DS-Qwen-1.5B	50 → 40	0.971	0.970	0.833 / 0.940	0.005
DS-Qwen-1.5B	50 → 60	0.977	0.977	0.851 / 0.956	0.005
DS-Qwen-1.5B	60 → 40	0.960	0.970	0.827 / 0.933	0.010
DS-Qwen-1.5B	60 → 50	0.970	0.976	0.834 / 0.944	0.005
DS-Qwen-7B	40 → 50	0.930	0.914	0.811 / 0.905	0.005
DS-Qwen-7B	40 → 60	0.916	0.922	0.803 / 0.899	0.010
DS-Qwen-7B	50 → 40	0.942	0.964	0.803 / 0.919	0.015
DS-Qwen-7B	50 → 60	0.917	0.922	0.796 / 0.903	0.010
DS-Qwen-7B	60 → 40	0.934	0.964	0.804 / 0.922	0.020
DS-Qwen-7B	60 → 50	0.919	0.914	0.793 / 0.903	0.020
DS-Qwen-14B	50 → 60	0.971	0.974	0.828 / 0.937	0.010
DS-Qwen-14B	60 → 50	0.956	0.975	0.799 / 0.920	0.015

(b) MATH-500

Model	Levels	Transfer	Own	Shuffled	p
DS-Qwen-1.5B	40 → 50	0.971	0.976	0.868 / 0.956	0.005
DS-Qwen-1.5B	40 → 60	0.979	0.985	0.881 / 0.964	0.005
DS-Qwen-1.5B	50 → 40	0.972	0.971	0.858 / 0.955	0.005
DS-Qwen-1.5B	50 → 60	0.985	0.985	0.879 / 0.964	0.005
DS-Qwen-1.5B	60 → 40	0.966	0.971	0.852 / 0.951	0.005
DS-Qwen-1.5B	60 → 50	0.972	0.976	0.865 / 0.959	0.005
DS-Qwen-7B	50 → 60	0.985	0.984	0.871 / 0.970	0.005
DS-Qwen-7B	50 → 70	0.988	0.991	0.907 / 0.981	0.015
DS-Qwen-7B	60 → 50	0.973	0.975	0.864 / 0.961	0.010
DS-Qwen-7B	60 → 70	0.991	0.991	0.888 / 0.979	0.005
DS-Qwen-7B	70 → 50	0.967	0.975	0.854 / 0.957	0.035
DS-Qwen-7B	70 → 60	0.978	0.984	0.852 / 0.963	0.015
DS-Qwen-14B	60 → 70	0.973	0.983	0.879 / 0.960	0.010
DS-Qwen-14B	70 → 60	0.980	0.985	0.868 / 0.954	0.005
OT3-7B	60 → 70	0.982	0.941	0.847 / 0.931	0.005
OT3-7B	70 → 60	0.733	0.939	0.686 / 0.784	0.204

Table 21: Held-out ROC AUC for separating terminating from non-terminating continuations when a direction constructed at one sparsity level (%) is applied to another. **Transfer**: AUC of the source direction on the target's held-out continuations; **Own**: AUC of the target's own direction; **Shuffled**: mean / 95th percentile of the transfer AUC when the source labels are permuted (200 permutations), with permutation p-value p.

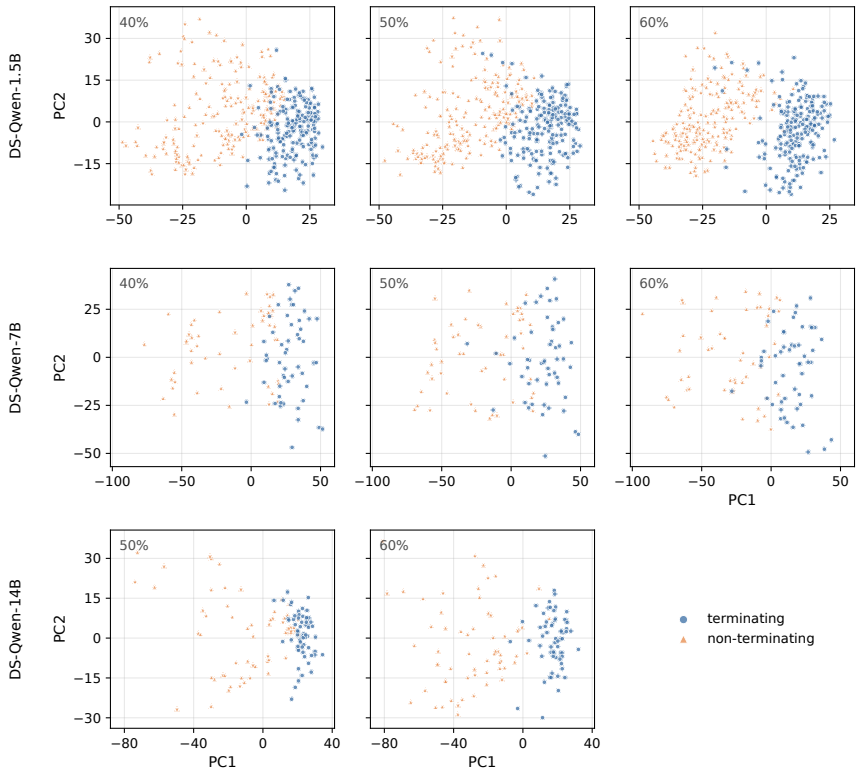

Figure 10: PCA of continuation representations at the steering layer on OpenR1-Math. For each model, PCA is fitted jointly on the continuations of all sparsity levels. Terminating continuations form a compact cluster, while non-terminating ones spread widely along the same PC1 at every sparsity level.

E Additional Main Experimental Results

E.1 Comprehensive Evaluation of Training-free Methods

Overall accuracy. Tables 23 and 24 show that TGAS achieves the best or tied-best accuracy on 53 of 72 checkpoint-benchmark pairs, compared with 10 for Forced-Answer. It ranks first in 20 of 24 model–sparsity settings and outperforms Base in 22, with the mean gain increasing from 1.6 points on dense models to 3.7 points at 60% sparsity. For OT3-7B at 60% sparsity, TGAS improves LiveCodeBench and AIME-25 accuracy by 5 points each. TGAS underperforms Base on only 8 of 72 pairs, five on AIME-25. TGAS improves termination over Base on 65 of 72 pairs, including 51 of 54 pruned settings, with gains increasing from 2.6 points on dense models to 11.2 points at 60% sparsity. These gains generally translate into higher accuracy: among the 65 pairs with improved termination, accuracy increases on 53 and decreases on only 6. For example, on MATH-500 at 60% sparsity, TGAS raises termination and accuracy from 21.9% to 67.9% and 15.5% to 30.9%, respectively, for DS-Qwen-1.5B.

Comparison with SEAL. SEAL improves termination by a similar margin, 5.8 points on average, compared with 5.9 for TGAS and raises it on 62 of 72 pairs, but often at the cost of accuracy. Its accuracy is below Base on 40 of 72 pairs, including 33 of the 62 pairs where it increases termination, and its average accuracy falls below Base in 14 of 24 model-sparsity settings. For example, on DS-Llama-8B at 10% sparsity, SEAL raises AIME-25 termination from 74.2% to 88.3% while accuracy drops from 35.8% to 30.8%; on DS-Qwen-7B at 60% sparsity, AIME-25 termination rises from 43.3% to 60.8% while accuracy drops from 22.5% to 16.7%. TGAS is more accurate than SEAL on 62 of 72 pairs and less accurate on only 3, all on AIME-25 with dense or 10%-sparsity OT3 models.

Forced-Answer and No-Repeat. Forced-Answer is most competitive at 60% sparsity, accounting for 7 of its 10 best or tied-best results and improving average accuracy by 2.4 points over Base. However, it reduces accuracy on 27 of 36 dense and 10%-sparsity pairs, limiting its practicality when problem difficulty is unknown. No-Repeat yields inconsistent gains, ranking best or tied-best on only 4 pairs while reducing accuracy on 41 of 72. It also underperforms Base on LiveCodeBench by 6.9 points on average, sometimes increasing termination at the expense of accuracy, e.g., 80.3% vs. 75.8% termination but 46.6% vs. 62.4% accuracy for Qwen-14B at 50% sparsity.

E.2 Comparison of Training-Free and Training-Based Methods

Table 25 summarizes the results. At 50% sparsity, training-free methods match or outperform Masked-LoRA without additional training. At 60% sparsity, Masked-LoRA achieves stronger recovery, improving LiveCodeBench accuracy from 14.7% to 33.4% for DS-Qwen-7B and from 37.6% to 51.8% for DS-Qwen-14B, surpassing the best training-free results (19.5% and 42.9%). On MATH-500, it reaches 37.8%, 74.8%, and 79.0% for Qwen-1.5B, 7B, and 14B, respectively. These results reveal a trade-off: training-based recovery is more effective under severe pruning but requires optimization and a separate LoRA adapter for each sparsity level, whereas TGAS remains a competitive alternative without weight updates.

E.3 Headroom and Recovery Analysis

Pruned models often fail to terminate even after their reasoning has already reached the correct answer. We ask two questions: how much accuracy would be regained if these generations simply committed to the answer they already state, and how many non-terminating generations TGAS actually turns into correct ones.

Setup. We evaluate DS-Qwen-1.5B, 7B, and 14B on 262 level 4-5 MATH-500 problems, from dense to 70% sparsity, using temperature 0.6 and a 16,384-token budget. We define oracle accuracy as Base accuracy plus the fraction of non-terminating generations that already state the gold answer *Conclusively* (the same rule in Appendix C.2). For recovery, we measure the fraction of initially incorrect non-terminating and terminated-wrong generations corrected by TGAS on the same problems and seeds, and the fraction of Base-correct generations that remain correct.

Headroom. At 50% sparsity, the oracle is at most 2.6 points above Base (Table 22). Headroom

			Accuracy (%)			Under TGAS (%)		
Model	**Ckpt**	**#NT**	**Base**	**Oracle**	**TGAS**	**NT→Correct**	**Wrong→Correct**	**Kept**
DS-Qwen-1.5B	dense	22	75.2±0.5	75.8±1.4	77.1±5.4	19.6±4.0	43.2±9.5	90.9±1.5
	obc_10	22	73.3±2.7	73.7±2.7	76.7±0.0	29.0±9.5	42.5±6.5	90.7±4.8
	obc_50	72	54.8±2.4	57.4±3.0	61.8±4.9	33.7±1.9	42.2±13.9	82.2±2.3
	obc_60	204	15.5±3.0	32.6±3.0	30.9±1.1	25.8±0.3	9.2±5.5	65.8±3.9
DS-Qwen-7B	dense	4	85.1±1.1	85.1±1.1	86.3±1.9	16.6±23.5	44.0±3.7	94.2±2.5
	obc_10	6	84.7±0.0	84.9±0.3	88.2±0.8	52.6±53.1	45.6±6.2	95.7±0.3
	obc_50	20	83.6±1.1	84.5±1.4	83.6±0.8	52.3±7.6	37.5±1.2	91.3±1.8
	obc_60	57	70.2±0.5	75.2±0.5	73.7±0.3	36.1±2.8	45.4±3.4	88.6±0.9
	obc_70	190	21.9±0.8	43.7±3.0	39.7±3.0	31.2±3.4	7.1±6.9	74.6±5.7
DS-Qwen-14B	dense	2	88.4±0.3	88.4±0.3	89.3±0.5	48.6±68.8	33.4±16.8	96.5±1.2
	obc_10	4	87.8±1.1	87.8±1.1	89.7±0.8	51.5±6.6	54.9±11.4	94.6±2.7
	obc_50	12	82.4±0.5	82.8±0.0	86.3±0.5	24.3±8.8	52.4±8.2	95.1±0.9
	obc_60	52	67.6±1.6	70.0±0.3	72.5±1.3	44.2±5.5	37.6±4.2	87.3±2.5
	obc_70	201	15.7±1.1	33.0±0.8	22.1±2.6	17.5±2.7	5.3±5.1	51.2±7.6

Table 22: Headroom and recovery on MATH-500. **#NT**: number of non-terminating Base generations. **Oracle**: Base accuracy if every non-terminating generation that already states the gold answer explicitly were counted as correct. **NT→Correct** and **Wrong→Correct**: share of non-terminating and terminated-wrong Base generations that become correct under TGAS. **Kept**: share of Base-correct generations that remain correct.

grows under heavier pruning: committing to an already stated answer would raise DS-Qwen-1.5B at 60% from 15.5% to 32.6%, and DS-Qwen-7B and 14B at 70% from 21.9% to 43.7% and from 15.7% to 33.0%. TGAS captures 90% and 70% of this headroom for DS-Qwen-1.5B and 7B at 60%, exceeds the oracle for DS-Qwen-14B, and captures 82% and 37% for DS-Qwen-7B and 14B at 70%.

Recovery. At 60% sparsity, TGAS turns 25.8%, 36.1% and 44.2% of non-terminating generations into correct ones for DS-Qwen-1.5B, 7B and 14B, against 17.1% (pooled over models) for a rerun of the unsteered model with another seed. In contrast, it does not help generations that terminated with a wrong answer: pooled over models, 31.2% of them become correct under TGAS and 31.9% under the rerun. TGAS therefore restores termination rather than correcting wrong answers, consistent with the intended behavior of its steering method.

Model	Ckpt	Method	MATH-500 Acc	MATH-500 Term	AIME-25 Acc	AIME-25 Term	LiveCodeBench Acc	LiveCodeBench Term
Deepseek-R1-Distill -Qwen-1.5B	dense	Base	75.2±0.5	91.6±0.0	22.5±4.2	71.7±5.8	34.5±1.6	76.1±1.8
		No-Repeat	71.0±0.0	91.2±0.5	22.5±4.2	74.2±7.4	27.6±1.6	80.8±3.9
		Forced-Answer	74.2±2.4	—	22.5±4.2	—	35.3±1.8	—
		SEAL	73.7±2.1	95.8±0.6	18.3±1.9	76.7±6.1	35.0±3.7	84.7±3.3
		TGAS	**77.1±5.4**	92.4±1.3	**25.8±4.2**	75.8±5.7	**36.8±2.6**	86.3±3.7
	obc_10	Base	73.3±2.7	91.4±1.3	21.7±4.3	71.7±1.9	32.9±2.6	76.6±3.1
		No-Repeat	74.0±0.5	91.6±1.6	24.2±5.7	70.0±8.6	30.5±2.3	82.4±3.0
		Forced-Answer	72.5±2.2	—	21.7±4.3	—	33.4±3.4	—
		SEAL	74.4±1.4	95.0±1.3	21.7±4.3	80.8±7.4	34.5±4.6	85.5±3.0
		TGAS	**76.7±0.0**	95.4±1.1	**26.7±4.7**	75.0±6.9	**37.1±2.6**	83.2±2.9
	obc_50	Base	54.8±2.4	72.5±0.5	13.3±2.7	46.7±4.7	13.4±2.2	32.1±4.3
		No-Repeat	57.8±1.9	80.3±1.9	**17.5±5.7**	51.7±7.9	9.5±2.6	40.8±6.2
		Forced-Answer	59.4±3.0	—	13.3±2.7	—	**16.6±3.1**	—
		SEAL	56.5±3.1	82.4±1.4	15.8±3.2	60.0±6.7	13.2±3.5	46.3±4.0
		TGAS	**61.8±4.9**	85.1±1.9	15.8±1.7	58.3±5.8	16.1±1.0	51.8±6.9
	obc_60	Base	15.5±3.0	21.9±1.9	1.7±1.9	6.7±4.7	0.0±0.0	0.5±0.6
		No-Repeat	26.3±7.0	42.4±5.4	2.5±5.0	23.3±4.7	0.0±0.0	7.4±0.9
		Forced-Answer	**34.7±3.8**	—	2.5±3.2	—	0.0±0.0	—
		SEAL	30.2±3.7	53.4±2.9	5.0±6.4	27.5±7.4	0.0±0.0	0.0±0.0
		TGAS	30.9±1.1	67.9±1.3	**5.0±1.9**	40.0±4.7	0.0±0.0	2.4±1.8
Deepseek-R1-Distill -Qwen-7B	dense	Base	85.1±1.1	98.3±0.3	**42.5±5.7**	80.0±7.2	59.2±3.8	83.7±1.4
		No-Repeat	83.8±0.8	96.8±0.8	34.2±5.7	77.5±5.0	45.0±1.8	82.4±3.4
		Forced-Answer	84.2±0.3	—	42.5±5.7	—	53.9±3.5	—
		SEAL	84.4±1.0	98.1±0.7	36.7±2.7	80.8±3.2	57.9±1.2	90.0±2.6
		TGAS	**86.3±1.9**	97.7±1.1	41.7±6.9	79.2±8.8	**59.7±2.2**	89.2±3.0
	obc_10	Base	84.7±0.0	97.7±1.1	31.7±5.8	80.0±7.2	58.4±1.8	82.6±2.6
		No-Repeat	84.2±0.8	96.6±1.1	35.0±3.3	78.3±4.3	49.5±2.3	84.5±1.6
		Forced-Answer	83.4±0.8	—	31.7±5.8	—	53.4±2.6	—
		SEAL	81.7±0.6	96.9±1.0	32.5±10.7	80.8±5.7	55.8±2.6	86.1±1.8
		TGAS	**88.2±0.8**	99.2±0.8	**43.3±4.7**	80.8±1.7	**60.3±3.4**	90.8±3.8
	obc_50	Base	83.6±1.1	92.2±1.3	30.0±4.7	57.5±7.4	**42.1±1.7**	63.9±3.1
		No-Repeat	80.5±1.6	91.6±2.2	31.7±8.8	65.8±3.2	33.7±2.7	66.1±5.0
		Forced-Answer	81.9±0.8	—	30.0±4.7	—	40.8±1.6	—
		SEAL	79.4±1.1	93.5±1.4	29.2±5.7	70.8±5.0	40.3±1.6	75.0±3.9
		TGAS	**83.6±0.8**	95.0±0.8	**31.7±3.3**	68.3±6.9	41.8±2.3	66.6±2.2
	obc_60	Base	70.2±0.5	78.2±0.5	22.5±6.3	43.3±8.6	14.7±4.3	34.2±5.0
		No-Repeat	69.3±1.3	80.9±1.6	22.5±4.2	48.3±4.3	14.7±1.5	50.0±2.5
		Forced-Answer	**75.0±0.3**	—	22.5±3.2	—	16.3±4.0	—
		SEAL	71.4±0.6	85.9±0.8	16.7±4.7	60.8±7.9	18.4±1.4	48.4±2.7
		TGAS	73.7±0.3	88.5±0.3	**22.5±3.2**	54.2±3.2	**19.5±3.3**	50.5±3.0
Deepseek-R1-Distill -Qwen-14B	dense	Base	88.4±0.3	99.0±0.3	54.2±5.0	87.5±5.0	73.2±3.9	90.8±2.2
		No-Repeat	87.6±0.8	98.1±1.1	45.8±11.3	80.8±8.3	56.8±1.5	89.5±1.9
		Forced-Answer	86.1±0.3	—	**55.0±5.8**	—	63.4±1.8	—
		SEAL	88.9±0.5	98.5±0.6	45.0±3.3	86.7±3.8	72.6±2.3	92.4±2.3
		TGAS	**89.3±0.5**	99.2±0.5	48.3±6.4	83.3±3.8	**73.4±1.8**	92.1±2.0
	obc_10	Base	87.8±1.1	98.5±0.5	47.5±6.9	84.2±5.0	71.8±2.3	89.2±2.3
		No-Repeat	89.5±0.3	98.9±0.5	47.5±10.0	83.3±2.7	56.1±1.8	88.9±2.2
		Forced-Answer	86.3±1.1	—	49.2±5.0	—	63.9±1.3	—
		SEAL	86.6±0.8	98.9±0.3	45.0±4.3	85.8±5.7	72.4±1.6	93.4±2.8
		TGAS	**89.7±0.8**	98.5±0.0	**50.8±3.2**	87.5±1.7	**72.9±2.3**	94.2±2.6
	obc_50	Base	82.4±0.5	95.4±0.5	35.8±5.0	63.3±6.1	62.4±1.0	75.8±2.3
		No-Repeat	83.6±1.1	95.0±1.6	37.5±3.2	64.2±6.3	46.6±2.9	80.3±3.6
		Forced-Answer	81.1±1.9	—	36.7±4.7	—	56.6±3.3	—
		SEAL	84.4±0.8	98.1±0.5	42.5±5.7	70.0±3.8	61.3±2.8	78.7±2.9
		TGAS	**86.3±0.5**	95.4±0.3	**42.5±3.2**	70.8±1.7	**62.9±3.5**	77.1±3.5
	obc_60	Base	67.6±1.6	80.2±0.5	26.7±4.7	39.2±5.0	37.6±5.6	46.6±3.9
		No-Repeat	71.6±1.3	91.4±0.8	25.0±3.3	70.0±10.5	33.9±1.0	60.3±3.0
		Forced-Answer	70.8±0.8	—	**27.5±3.2**	—	38.9±4.3	—
		SEAL	71.0±0.8	86.6±1.8	24.2±4.2	45.0±6.4	42.9±2.9	60.0±0.9
		TGAS	**72.5±1.3**	88.9±2.2	25.8±3.2	44.2±4.2	**42.9±2.6**	53.7±3.1

Table 23: Training-free remedies for SparseGPT-pruned LRMs with OBC calibration DeepSeek-R1-Distill-Qwen models. Acc. and Term. denote accuracy and termination rate, respectively. Forced-Answer's Term. is not reported, since forcing ends every generation by design.

Model	Ckpt	Method	MATH-500		AIME-25		LiveCodeBench	
			Acc	**Term**	**Acc**	**Term**	**Acc**	**Term**
Deepseek-R1-Distill -Llama-8B	dense	Base	82.4±2.7	97.9±0.8	34.2±5.7	70.0±11.2	58.7±2.5	85.5±1.8
		No-Repeat	83.0±0.8	97.7±0.0	35.0±4.3	71.7±3.3	44.7±3.3	85.0±1.3
		Forced-Answer	81.3±1.6	—	30.8±1.7	—	53.2±1.1	—
		SEAL	80.2±1.8	97.7±0.3	32.5±5.7	87.5±4.2	57.9±2.7	91.6±2.3
		TGAS	**84.0±2.2**	97.3±0.0	**38.3±5.8**	80.8±3.2	**60.8±2.2**	87.9±1.8
	obc_10	Base	**84.0±3.8**	96.4±1.9	35.8±6.9	74.2±1.7	**60.3±2.2**	87.1±0.5
		No-Repeat	82.3±1.9	96.4±0.8	35.8±8.8	75.0±4.3	43.9±4.1	86.6±3.1
		Forced-Answer	81.5±2.4	—	30.0±2.7	—	53.2±1.1	—
		SEAL	80.5±2.7	98.1±0.9	30.8±5.0	88.3±6.9	55.5±2.2	93.2±1.1
		TGAS	83.6±2.7	97.3±0.0	**35.8±5.7**	79.2±11.0	59.5±3.7	90.3±4.3
	obc_50	Base	65.5±0.3	88.5±2.2	21.7±7.9	75.0±8.4	42.6±0.6	66.8±1.8
		No-Repeat	64.7±0.8	90.6±1.3	20.8±8.8	76.7±2.7	33.2±3.0	75.3±4.9
		Forced-Answer	64.5±0.5	—	16.7±6.1	—	40.8±1.3	—
		SEAL	63.0±0.6	89.3±1.1	23.3±5.4	76.7±9.4	41.3±4.1	75.3±3.9
		TGAS	**67.6±1.6**	88.5±1.1	**23.3±3.8**	80.8±4.2	**43.9±3.0**	72.1±2.0
	obc_60	Base	30.0±3.0	50.6±2.4	3.3±2.7	29.2±4.2	10.8±1.6	26.3±2.3
		No-Repeat	33.4±0.8	74.4±1.1	2.5±3.2	50.8±10.7	9.5±1.5	42.1±3.7
		Forced-Answer	35.7±1.3	—	3.3±2.7	—	**11.1±1.4**	—
		SEAL	32.8±1.6	61.8±1.3	4.2±5.0	45.0±4.3	10.0±3.0	32.6±1.9
		TGAS	**36.6±1.6**	65.6±0.5	**7.5±1.7**	45.0±4.3	10.8±1.8	33.7±1.5
OpenThinker3 -1.5B	dense	Base	80.2±1.6	87.2±1.9	32.5±1.7	39.2±3.2	43.2±2.9	57.4±3.3
		No-Repeat	80.2±2.2	87.4±1.6	35.0±5.8	44.2±4.2	42.1±2.3	58.2±2.3
		Forced-Answer	79.2±3.0	—	30.8±3.2	—	41.1±2.4	—
		SEAL	80.2±2.0	89.7±1.3	**37.5±5.0**	45.8±1.7	41.6±2.0	61.3±3.5
		TGAS	**84.0±1.1**	89.7±0.5	36.7±2.7	43.3±3.8	**45.3±2.3**	61.1±0.9
	obc_10	Base	79.6±1.3	87.0±0.5	34.2±8.3	40.0±2.7	43.4±2.3	57.1±2.6
		No-Repeat	79.2±1.9	86.6±2.2	35.0±6.9	41.7±6.9	38.9±2.6	55.5±1.6
		Forced-Answer	77.3±3.0	—	29.2±6.3	—	40.8±2.6	—
		SEAL	77.1±1.9	87.0±2.0	**41.7±5.8**	52.5±5.7	38.7±2.0	61.3±5.8
		TGAS	**81.7±1.9**	88.5±0.8	37.5±3.2	46.7±9.8	**45.5±1.6**	61.3±2.8
	obc_50	Base	56.9±3.8	67.4±4.0	**20.8±1.7**	32.5±5.0	10.0±0.6	22.4±1.3
		No-Repeat	54.4±0.3	67.4±1.3	16.7±4.7	25.8±3.2	6.3±0.9	23.2±4.5
		Forced-Answer	57.1±3.5	—	17.5±5.0	—	11.1±1.8	—
		SEAL	55.0±5.1	73.7±2.1	18.3±4.3	34.2±1.7	9.2±1.0	26.8±2.8
		TGAS	**58.4±2.2**	74.4±0.3	19.2±7.4	44.2±5.7	**13.9±2.3**	29.2±5.6
	obc_60	Base	17.6±1.1	13.4±2.2	0.8±1.7	5.8±1.7	0.0±0.0	0.0±0.0
		No-Repeat	19.1±1.1	15.8±1.9	2.5±1.7	7.5±3.2	0.0±0.0	0.3±0.5
		Forced-Answer	**23.1±0.3**	—	**3.3±0.0**	—	0.0±0.0	—
		SEAL	14.5±1.2	8.4±1.6	0.8±1.7	5.8±6.9	0.0±0.0	0.5±0.6
		TGAS	22.5±2.4	20.6±0.8	**3.3±0.0**	10.8±3.2	0.0±0.0	2.6±0.6
OpenThinker3 -7B	dense	Base	80.3±4.6	92.0±1.6	44.2±5.0	50.0±4.7	63.4±2.0	72.9±1.8
		No-Repeat	80.9±1.1	91.4±0.8	50.0±8.2	55.8±10.0	58.2±1.6	71.6±2.7
		Forced-Answer	80.2±2.7	—	32.5±1.7	—	53.7±0.9	—
		SEAL	79.4±2.5	93.5±1.2	**50.8±6.9**	53.3±7.2	63.4±1.3	75.3±1.8
		TGAS	**84.4±1.6**	93.5±1.3	45.8±4.2	54.2±3.2	**64.7±1.4**	73.9±1.8
	obc_10	Base	81.7±3.8	90.8±0.5	**50.0±0.0**	55.0±1.9	61.3±2.2	71.8±1.3
		No-Repeat	80.7±3.5	92.0±1.1	46.7±2.7	53.3±4.7	56.3±1.8	68.9±3.9
		Forced-Answer	82.3±5.1	—	35.8±3.2	—	52.4±2.2	—
		SEAL	83.6±3.2	94.7±0.7	46.7±4.7	60.0±7.2	62.1±1.7	74.5±1.8
		TGAS	**84.7±0.3**	94.3±0.5	48.3±5.8	56.7±5.4	**63.4±0.5**	74.2±1.8
	obc_50	Base	84.5±1.9	89.7±2.2	37.5±4.2	50.8±6.9	51.1±1.1	63.2±2.6
		No-Repeat	80.9±0.5	88.5±1.1	39.2±5.7	50.8±3.2	47.6±4.0	65.8±3.0
		Forced-Answer	82.3±0.3	—	28.3±4.3	—	47.4±1.7	—
		SEAL	81.7±1.1	88.9±1.2	35.0±1.9	54.2±5.0	51.1±3.6	68.7±2.6
		TGAS	**85.5±1.6**	91.6±0.3	**43.3±8.2**	52.5±9.6	**55.5±2.2**	69.5±3.0
	obc_60	Base	72.3±1.3	83.6±1.1	24.2±5.7	47.5±6.3	31.3±2.6	54.2±1.8
		No-Repeat	71.2±0.3	83.2±0.5	25.8±4.2	44.2±5.7	26.3±3.0	53.9±4.4
		Forced-Answer	71.9±0.8	—	20.8±3.2	—	32.4±1.8	—
		SEAL	71.0±0.9	87.0±0.6	27.5±4.2	50.8±4.2	30.3±3.7	60.5±2.2
		TGAS	**74.8±1.1**	84.4±0.3	**29.2±1.7**	50.8±6.9	**36.3±3.0**	58.9±3.1

Table 24: Training-free remedies for SparseGPT-pruned LRMs with OBC calibration DeepSeek-R1-Distill-Llama and OpenThinker3 models. Acc. and Term. denote accuracy and termination rate, respectively. Forced-Answer's Term. is not reported, since forcing ends every generation by design.

Model	Ckpt	Method	MATH-500		AIME-25		LiveCodeBench	
			Acc	**Term**	**Acc**	**Term**	**Acc**	**Term**
Deepseek-R1-Distill -Qwen-1.5B	obc_50	Base	54.8±2.4	72.5±0.5	13.3±2.7	46.7±4.7	13.4±2.2	32.1±4.3
		Forced-Answer	59.4±3.0	—	13.3±2.7	—	**16.6±3.1**	—
		TGAS	**61.8±4.9**	85.1±1.9	**15.8±1.7**	58.3±5.8	16.1±1.0	51.8±6.9
		Masked-LoRA	60.3±3.1	82.1±1.4	11.7±5.8	52.5±5.7	15.3±1.8	30.8±4.3
	obc_60	Base	15.5±3.0	21.9±1.9	1.7±1.9	6.7±4.7	0.0±0.0	0.5±0.6
		Forced-Answer	34.7±3.8	—	2.5±3.2	—	0.0±0.0	—
		TGAS	30.9±1.1	67.9±1.3	5.0±1.9	40.0±4.7	0.0±0.0	2.4±1.8
		Masked-LoRA	**37.8±3.7**	62.2±2.9	**5.8±4.2**	45.8±5.7	**0.3±0.5**	5.3±1.5
Deepseek-R1-Distill -Qwen-7B	obc_50	Base	83.6±1.1	92.2±1.3	30.0±4.7	57.5±7.4	**42.1±1.7**	63.9±3.1
		Forced-Answer	81.9±0.8	—	30.0±4.7	—	40.8±1.6	—
		TGAS	**83.6±0.8**	95.0±0.8	**31.7±3.3**	68.3±6.9	41.8±2.3	66.6±2.2
		Masked-LoRA	81.7±1.1	89.3±1.4	29.2±3.2	52.5±4.2	41.6±1.4	60.0±2.3
	obc_60	Base	70.2±0.5	78.2±0.5	22.5±6.3	43.3±8.6	14.7±4.3	34.2±5.0
		Forced-Answer	**75.0±0.3**	—	22.5±3.2	—	16.3±4.0	—
		TGAS	73.7±0.3	88.5±0.3	22.5±3.2	54.2±3.2	19.5±3.3	50.5±3.0
		Masked-LoRA	74.8±0.6	85.5±0.8	**26.7±2.7**	51.7±7.9	**33.4±2.3**	52.6±3.3
	obc_70	Base	21.9±0.8	24.8±0.5	5.0±3.3	19.2±3.2	0.0±0.0	0.0±0.0
		Forced-Answer	43.3±1.3	—	5.0±3.3	—	0.0±0.0	—
		TGAS	39.7±1.1	61.1±0.5	5.8±4.2	23.3±8.6	0.0±0.0	5.0±1.0
		Masked-LoRA	**56.1±1.1**	76.3±0.5	**15.0±3.3**	45.8±5.0	**0.5±1.1**	20.0±3.0
Deepseek-R1-Distill -Qwen-14B	obc_50	Base	82.4±0.5	95.4±0.5	35.8±5.0	63.3±6.1	62.4±1.0	75.8±2.3
		Forced-Answer	81.1±1.9	—	36.7±4.7	—	56.6±3.3	—
		TGAS	**86.3±0.5**	95.4±0.3	**42.5±3.2**	70.8±1.7	**62.9±3.5**	77.1±3.5
		Masked-LoRA	82.1±1.0	92.4±0.8	32.5±5.0	49.2±1.7	60.0±0.9	69.7±2.3
	obc_60	Base	67.6±1.6	80.2±0.5	26.7±4.7	39.2±5.0	37.6±5.6	46.6±3.9
		Forced-Answer	70.8±0.8	—	**27.5±3.2**	—	38.9±4.3	—
		TGAS	72.5±1.3	88.9±2.2	25.8±3.2	44.2±4.2	42.9±2.6	53.7±3.1
		Masked-LoRA	**79.0±1.3**	88.9±1.2	24.2±5.0	46.7±7.2	**51.8±1.6**	65.0±3.4
	obc_70	Base	15.6±1.1	23.1±0.3	3.3±2.7	10.8±3.2	0.0±0.0	0.0±0.0
		Forced-Answer	31.9±1.3	—	5.8±4.2	—	0.0±0.0	—
		TGAS	22.1±1.2	37.4±0.3	5.8±1.7	16.7±2.7	0.0±0.0	0.3±0.5
		Masked-LoRA	**58.0±1.2**	72.5±0.3	**12.5±1.7**	34.2±5.0	**17.9±3.3**	33.2±2.5

Table 25: Training-free remedies vs. Masked-LoRA for SparseGPT-pruned LRMs with OBC calibration DeepSeek-R1-Distill-Qwen models. Forced-Answer's **Term** is not reported, since forcing ends every generation by design.

Model	Ckpt	Δ_{cyc} at decisions U	Δ_{cyc} at decisions S	DPs/1k U	DPs/1k S
Deepseek-R1-Distill -Qwen-1.5B	dense	−14.9±0.5	−16.4±0.6	8.1	5.6
	obc_10	−15.3±0.5	−16.5±0.6	8.0	5.3
	obc_50	−10.2±0.4	−12.7±0.5	9.2	5.2
	obc_60	−5.6±0.3	−6.3±0.4	23.7	10.5
Deepseek-R1-Distill -Qwen-7B	dense	−18.2±0.6	−20.5±0.6	7.1	4.7
	obc_10	−18.2±0.5	−19.5±0.6	7.4	4.5
	obc_50	−15.3±0.5	−17.1±0.5	9.0	5.4
	obc_60	−12.0±0.4	−15.4±0.5	12.3	5.8
Deepseek-R1-Distill -Qwen-14B	dense	−16.1±0.5	−18.1±0.5	6.6	4.8
	obc_10	−16.0±0.4	−18.2±0.5	6.7	4.8
	obc_50	−14.7±0.4	−17.9±0.5	7.2	5.5
	obc_60	−12.7±0.4	−14.0±0.4	11.5	6.8

Table 26: Δ_{cyc} at decision points on MATH-500 L4–5 with and without TGAS. **U**: unsteered traces; **S**: steered traces scored with the steering hook. More negative values indicate stronger preference for moving on. **DPs/1k**: decision points per 1,000 generated tokens.

F Additional Ablation Studies

F.1 Behavioral Changes in Pruned Models After Steering

Section 3.2 showed that pruning makes re-treading increasingly competitive at decision points. $\Delta_{\text{cyc}} = \log P(\text{CYCLIC}) - \log P(\text{PROGRESS})$, measured over the next 32 tokens after a reconsideration cue (*"Wait"*, *"Alternatively"*, ...), approaches zero as sparsity increases. We ask whether TGAS reverses this trend. We repeat the measurement on steered DeepSeek-R1-Distill-Qwen models, using decision points from the steered traces.

Table 26 reports MATH-500 results for the three DeepSeek-R1-Distill-Qwen models. TGAS makes Δ_{cyc} more negative in all 12 settings, with shifts of 0.7-3.4. More consistently, it reduces decision points by 24%-56%, with the largest reductions at 60% sparsity: from 23.7 to 10.5, 12.3 to 5.8, and 11.5 to 6.8 per 1k tokens for the 1.5B, 7B, and 14B models, respectively. For the latter two, this restores reconsideration frequency to approximately dense-model levels. Overall, TGAS primarily suppresses repeated reconsideration and promotes earlier termination.

F.2 Ablation Study of Forced-Answer

Table 27 shows that the effectiveness of Forced-Answer depends strongly on the generation budget. Short budgets can truncate useful reasoning in capable models, e.g., dense Qwen-14B drops from 54.2% to 25.0% on AIME-25 at $D = 4{,}096$, while longer budgets are less effective for heavily pruned models. At $D = 16{,}384$, forcing only at the token limit is generally safe but provides limited gains. Overall, no single budget works across models and sparsity levels. In practice, selecting D requires estimating the difficulty of each problem, since a budget sufficient for difficult problems may unnecessarily truncate easier ones. This dependence on problem difficulty is a major practical drawback of Forced-Answer.

F.3 Ablation Study of No-Repeat

Table 28 shows that No-Repeat is sensitive to the blocked n-gram size and that improved termination does not necessarily translate into accuracy gains. With $n = 8$, accuracy drops across nearly all settings, including dense models, with particularly severe degradation on LiveCodeBench: DeepSeek-R1-Distill-Qwen-14B falls from 73.2% to 19.5%, and the 7B model from 59.2% to 15.0%. Although $n = 8$ improves termination for heavily pruned 1.5B models and on LiveCodeBench, accuracy remains near zero in these settings, indicating that most additional terminations are incorrect. It also reduces termination for the 7B model on AIME-25 at every pruned checkpoint, from 43.3% to 35.8% at 60% sparsity.

These results suggest that blocking short repeated spans can disrupt valid reasoning and code patterns. Increasing the block size to $n = 32$ mitigates accuracy degradation and preserves AIME-25 performance for most checkpoints, while improving termination at moderate sparsity, e.g., 39.2% to 70.0% for the 14B model at 60% sparsity. However, these gains yield limited accuracy recovery, and substantial degradation persists on LiveCodeBench, including for dense models. Overall, No-Repeat can suppress repetitive generation but does not reliably guide the model toward correct answers, making it the least effective of the three remedies.

F.4 Ablation Study of Steering Strength

Table 29 shows that the steering strength α provides direct control over termination, with stronger effects under higher sparsity. For DeepSeek-R1-Distill-Qwen-7B at 60% sparsity, LiveCodeBench termination increases from 1.8% at $\alpha = -2$ to 34.2% without steering and 57.9% at $\alpha = +2$. On AIME-25, it similarly increases from 13.3% to 43.3% and 55.0%, respectively. The effect is smaller for dense and lightly pruned check-

Model	Ckpt	AIME-25				LiveCodeBench		
		Base	D=4096	D=8192	D=16384	Base	D=4096	D=8192
Deepseek-R1-Distill -Qwen-1.5B	dense	22.5±4.2	19.2±1.7	22.5±4.2	22.5±4.2	34.5±1.6	35.0±1.8	35.3±1.8
	obc_10	21.7±4.3	16.7±4.7	21.7±6.9	21.7±4.3	32.9±2.6	33.7±1.9	33.4±3.4
	obc_50	13.3±2.7	12.5±3.2	14.2±1.7	13.3±2.7	13.4±2.2	19.2±4.6	16.6±3.1
	obc_60	1.7±1.9	2.5±3.2	2.5±3.2	2.5±3.2	0.0±0.0	0.0±0.0	0.0±0.0
	obc_70	0.0±0.0	0.0±0.0	0.0±0.0	0.0±0.0	0.0±0.0	0.0±0.0	0.0±0.0
Deepseek-R1-Distill -Qwen-7B	dense	42.5±5.7	25.8±3.2	39.2±4.2	42.5±5.7	59.2±3.8	50.3±2.2	53.9±3.5
	obc_10	31.7±5.8	24.2±4.2	30.8±3.2	31.7±5.8	58.4±1.8	50.0±2.0	53.4±2.6
	obc_50	30.0±4.7	23.3±2.7	25.8±5.0	30.0±4.7	42.1±1.7	40.5±1.8	40.8±1.6
	obc_60	22.5±3.2	15.8±4.2	20.8±1.7	22.5±3.2	14.7±4.3	15.8±3.1	16.3±4.0
	obc_70	5.0±3.3	6.7±0.0	5.8±1.7	5.0±3.3	0.0±0.0	0.0±0.0	0.0±0.0
Deepseek-R1-Distill -Qwen-14B	dense	54.2±5.0	25.0±5.8	44.2±3.2	55.0±5.8	73.2±3.9	52.9±1.0	63.4±1.8
	obc_10	47.5±6.9	26.7±2.7	39.2±3.2	49.2±5.0	71.8±2.3	53.9±3.0	63.9±1.3
	obc_50	35.8±5.0	26.7±2.7	33.3±4.7	36.7±4.7	62.4±1.0	50.0±4.7	56.6±3.3
	obc_60	26.7±4.7	20.8±3.2	25.0±3.3	27.5±3.2	37.6±5.6	30.5±1.5	38.9±4.3
	obc_70	3.3±2.7	5.0±4.3	5.8±4.2	5.8±4.2	0.0±0.0	0.0±0.0	0.0±0.0

Table 27: Effect of the Forced-Answer budget. Accuracy (%) after forcing an answer at D reasoning tokens, compared with the unforced checkpoint (*Base*). Generations finishing before D are kept unchanged, and at $D = 16{,}384$ only capped generations are forced.

		AIME-25						LiveCodeBench					
		Base		n=8		n=32		Base		n=8		n=32	
Model	**Ckpt**	**Acc**	**Term**	**Acc**	**Term**	**Acc**	**Term**	**Acc**	**Term**	**Acc**	**Term**	**Acc**	**Term**
DS-Qwen -1.5B	dense	22.5±4.2	71.7±5.8	7.5±4.2	68.3±7.9	22.5±4.2	74.2±7.4	34.5±1.6	76.1±1.8	8.9±1.4	79.7±1.0	27.6±1.6	80.8±3.9
	obc_10	21.7±4.3	71.7±1.9	6.7±2.7	66.7±4.7	24.2±5.7	70.0±8.6	32.9±2.6	76.6±3.1	8.9±2.6	77.9±4.8	30.5±2.3	82.4±3.0
	obc_50	13.3±2.7	46.7±4.7	3.3±0.0	46.7±9.8	17.5±5.7	51.7±7.9	13.4±2.2	32.1±4.3	1.8±1.0	40.3±5.0	9.5±2.6	40.8±6.2
	obc_60	1.7±1.9	6.7±4.7	0.0±0.0	42.5±13.4	2.5±5.0	23.3±4.7	0.0±0.0	0.5±0.6	0.0±0.0	31.6±3.6	0.0±0.0	7.4±0.9
DS-Qwen -7B	dense	42.5±5.7	80.0±7.2	23.3±7.2	76.7±8.2	34.2±5.7	77.5±5.0	59.2±3.8	83.7±1.4	15.0±1.6	82.1±2.3	45.0±1.8	82.4±3.4
	obc_10	31.7±5.8	80.0±7.2	29.2±4.2	79.2±3.2	35.0±3.3	78.3±4.3	58.4±1.8	82.6±2.6	17.1±4.1	84.5±2.8	49.5±2.3	84.5±1.6
	obc_50	30.0±4.7	57.5±7.4	20.8±3.2	48.3±6.9	31.7±8.8	65.8±3.2	42.1±1.7	63.9±3.1	10.0±2.2	67.6±3.1	33.7±2.7	66.1±5.0
	obc_60	22.5±3.2	43.3±8.6	12.5±5.0	35.8±3.2	22.5±4.2	48.3±4.3	14.7±4.3	34.2±5.0	4.2±1.7	40.5±5.0	14.7±1.5	50.0±2.5
	obc_70	5.0±3.3	19.2±3.2	0.0±0.0	2.5±3.2	4.2±1.7	15.0±3.3	0.0±0.0	0.0±0.0	0.0±0.0	3.9±1.6	0.0±0.0	1.6±1.4
DS-Qwen -14B	dense	54.2±5.0	87.5±5.0	33.3±4.7	82.5±5.7	45.8±11.3	80.8±8.3	73.2±3.9	90.8±2.2	19.5±4.4	93.2±2.5	56.8±1.5	89.5±1.9
	obc_10	47.5±6.9	84.2±5.0	40.8±6.3	87.5±7.4	47.5±10.0	83.3±2.7	71.8±2.3	89.2±2.3	18.7±3.7	94.2±1.4	56.1±1.8	88.9±2.2
	obc_50	35.8±5.0	63.3±6.1	28.3±4.3	66.7±11.9	37.5±3.2	64.2±6.3	62.4±1.0	75.8±2.3	11.6±3.3	83.9±2.8	46.6±2.9	80.3±3.6
	obc_60	26.7±4.7	39.2±5.0	17.5±1.7	45.0±5.8	25.0±3.3	70.0±10.5	37.6±5.6	46.6±3.9	9.5±1.9	63.9±4.6	33.9±1.0	60.3±3.0
	obc_70	3.3±2.7	10.8±3.2	1.7±1.9	10.0±4.7	5.0±5.8	24.2±10.7	0.0±0.0	0.0±0.0	0.0±0.0	7.1±2.3	0.0±0.0	1.8±1.3

Table 28: No-Repeat n-gram size. **Acc** is accuracy (%) and **Term** the share of generations that terminated (%). All three columns of a dataset come from the same evaluation pass, so *Base* here is that pass's unblocked run.

points, which already terminate reliably. Reversing the steering direction is particularly harmful: at 60% sparsity, LiveCodeBench accuracy drops from 14.7% to 1.1% when α changes from 0 to -2. These results indicate that the steering direction directly affects the model's stopping behavior.

However, increasing α too strong beyond zero primarily improves termination rather than accuracy. For DeepSeek-R1-Distill-Qwen-7B, positive α values produce little change in LiveCodeBench accuracy and even reduce the accuracy of AIME-25 at 60% sparsity: increasing the value of α from 1 to 2 reduces accuracy from 22.5% to 20%. OpenThinker3-7B exhibits a similar but weaker trend, with termination increasing modestly while accuracy changes only slightly. Overall, α is an effective control for termination, stronger steering mainly improves stopping behavior rather than solution accuracy, while reversing the steering direction can substantially degrade both.

F.5 Ablation on Number of Sampled Created Directions

Table 30 shows that a small number of samples is sufficient to obtain an effective steering direction. Using 50 samples per class achieves comparable steering performance to 100 samples or the full set for both DeepSeek-R1-Distill-Qwen-7B and OpenThinker3-7B. In particular, the termination rate remains stable across sample sizes: for DeepSeek-R1-Distill-Qwen-7B at 60% sparsity on LiveCodeBench, it stays within 52.9%-57.6%, compared with 34.2% without steering. These results suggest that 50 samples per class provide sufficient signal to identify an effective steering direction.

Model	Test	Ckpt	$\alpha=-2$		$\alpha=-1.5$		$\alpha=-1$		$\alpha=0$		$\alpha=+1$		$\alpha=+1.5$		$\alpha=+2$	
			Acc	**Term**	**Acc**	**Term**	**Acc**	**Term**	**Acc**	**Term**	**Acc**	**Term**	**Acc**	**Term**	**Acc**	**Term**
Deepseek-R1-Distill -Qwen-7B	AIME	dense	31.7±4.3	70.0±6.1	37.5±5.0	83.3±7.2	37.5±4.2	81.7±1.9	42.5±5.7	80.0±7.2	41.7±6.9	79.2±8.8	42.5±6.9	82.5±3.2	36.7±4.2	83.3±5.0
		obc_10	32.5±3.2	70.0±4.7	37.5±9.2	76.7±3.8	37.5±4.2	80.0±6.1	31.7±5.8	80.0±7.2	40.0±6.1	76.7±8.2	43.3±3.2	85.0±9.6	36.7±2.7	88.3±7.2
		obc_50	27.5±6.3	33.3±7.2	25.8±5.7	45.8±10.7	30.8±6.9	55.0±4.3	30.0±4.7	57.5±7.4	30.0±4.7	63.3±4.7	25.8±4.2	64.2±4.2	25.8±3.2	72.5±4.3
		obc_60	10.0±2.7	13.3±4.7	14.2±4.2	16.7±4.7	15.8±5.7	25.8±3.2	22.5±6.3	43.3±8.6	22.5±1.7	49.2±4.2	20.8±7.4	55.0±3.3	20.0±3.2	63.3±8.2
	LCB	dense	55.3±2.2	73.4±2.0	57.4±2.2	80.5±3.7	55.8±2.7	81.1±1.5	59.2±3.8	83.7±1.4	59.7±2.2	89.2±3.0	59.7±1.3	90.0±3.3	56.1±1.4	91.3±1.3
		obc_10	56.1±1.8	75.5±3.6	55.8±1.5	78.7±2.3	56.3±1.1	78.7±3.8	58.4±1.8	82.6±2.6	58.7±2.5	86.6±1.6	57.6±2.3	89.5±2.3	56.8±1.8	92.6±1.3
		obc_50	27.1±4.0	35.0±4.2	32.1±2.5	44.7±2.0	37.6±2.6	53.7±2.3	42.1±1.7	63.9±3.1	38.2±4.3	67.4±2.3	42.1±0.9	76.3±3.3	37.1±3.5	80.8±1.8
		obc_60	1.1±0.9	1.8±0.5	3.7±1.1	7.1±0.5	8.2±2.8	14.7±0.0	14.7±4.3	34.2±5.0	20.9±2.2	49.5±3.1	16.8±3.7	53.7±4.5	16.3±2.3	57.9±2.7
OpenThinker3 -7B	AIME	dense	47.5±5.7	49.2±4.2	45.8±6.9	49.2±5.0	42.5±1.7	45.8±4.2	44.2±5.0	50.0±4.7	42.5±4.2	47.5±3.2	45.8±4.2	54.2±3.2	43.3±5.4	50.0±5.4
		obc_10	41.7±4.3	43.3±4.7	47.5±3.2	48.3±4.3	42.5±4.2	47.5±6.3	50.0±0.0	55.0±1.9	48.3±7.9	55.0±6.9	48.3±5.8	56.7±5.4	45.8±4.2	52.5±5.0
		obc_50	35.0±3.3	40.8±3.2	30.0±9.4	41.7±10.4	31.7±4.3	40.0±6.1	37.5±4.2	50.8±6.9	37.5±6.3	48.3±7.9	39.2±7.4	53.3±8.6	39.2±10.0	50.0±7.2
		obc_60	19.2±7.4	35.8±8.8	24.2±4.2	42.5±3.2	23.3±3.8	39.2±6.9	24.2±5.7	47.5±6.3	24.2±5.0	50.0±10.5	29.2±1.7	50.8±6.9	22.5±5.0	51.7±6.4
	LCB	dense	60.8±2.3	69.5±0.9	61.3±2.3	71.3±3.4	61.6±3.0	71.8±4.4	63.4±2.0	72.9±1.8	62.1±1.5	74.2±0.6	64.7±1.4	73.9±1.8	61.3±2.3	72.1±2.8
		obc_10	61.8±1.3	68.9±1.4	61.6±1.4	70.8±3.6	61.8±1.0	69.7±2.2	61.3±2.2	71.8±1.3	60.5±0.6	71.1±1.1	63.4±0.5	74.2±1.8	63.2±1.9	75.0±3.4
		obc_50	52.6±0.9	64.7±2.2	52.4±2.3	63.7±5.3	53.2±2.2	64.2±2.3	51.1±1.1	63.2±2.6	51.8±1.8	66.8±1.4	52.4±3.5	67.1±3.7	53.4±1.0	67.6±2.3
		obc_60	24.7±2.5	49.5±2.1	28.4±2.3	54.2±3.9	30.0±1.4	54.7±2.9	31.3±2.6	54.2±1.8	32.6±3.1	58.9±3.7	36.3±3.0	58.9±3.1	38.2±2.2	60.0±3.5

Table 29: Effect of steering strength α on accuracy and termination. Each model is steered with its main direction scaled by α, absolute steering coefficient; $\alpha = 0$ is the unsteered checkpoint and negative α pushes against the direction. **Acc** is accuracy (%) and **Term** the share of generations that terminated (%).

Model	Test	Ckpt	n50		n75		n100		Full	
			Acc	**Term**	**Acc**	**Term**	**Acc**	**Term**	**Acc**	**Term**
Deepseek-R1-Distill -Qwen-7B	AIME	dense	38.3±6.4	82.5±9.2	36.7±9.0	79.2±5.7	42.5±4.2	80.0±3.8	41.7±6.9	79.2±8.8
		obc_10	35.8±3.2	83.3±8.6	38.3±4.3	80.8±3.2	35.8±1.7	80.8±3.2	43.3±4.7	80.8±1.7
		obc_50	31.7±4.3	69.2±3.2	28.3±6.4	62.5±5.0	30.8±8.8	69.2±4.2	31.7±3.3	68.3±6.9
		obc_60	18.3±4.3	55.0±5.8	20.0±3.8	57.5±3.2	25.0±6.4	62.5±3.2	22.5±3.2	54.2±3.2
	LCB	dense	58.2±2.3	87.4±3.1	57.9±2.3	86.3±0.9	57.9±2.3	88.4±2.6	59.7±2.2	89.2±3.0
		obc_10	57.9±1.5	87.1±4.1	53.2±2.0	87.9±3.0	56.6±2.0	89.7±1.8	60.3±3.4	90.8±3.8
		obc_50	38.9±5.8	73.2±3.9	40.8±2.3	74.5±1.8	37.9±2.6	73.9±3.8	41.8±2.3	66.6±2.2
		obc_60	16.3±2.2	57.6±3.1	16.1±2.2	53.9±5.0	16.3±1.4	52.9±3.0	19.5±3.3	50.5±3.0
OpenThinker3 -7B	AIME	dense	41.7±3.3	49.2±3.2	44.2±6.3	50.0±4.7	43.3±4.7	48.3±1.9	45.8±4.2	54.2±3.2
		obc_10	45.0±6.4	51.7±6.9	42.5±5.7	47.5±1.7	43.3±7.2	56.7±5.4	48.3±5.8	56.7±5.4
		obc_50	37.5±1.7	50.8±1.7	39.2±5.0	55.8±10.0	35.0±3.3	46.7±2.7	43.3±8.2	52.5±9.6
		obc_60	25.8±5.7	54.2±4.2	23.3±3.8	49.2±10.0	25.8±4.2	52.5±3.2	29.2±1.7	50.8±6.9
	LCB	dense	63.7±1.1	73.7±3.3	61.6±2.0	72.9±3.5	62.4±1.0	73.9±3.3	64.7±1.4	73.9±1.8
		obc_10	62.4±5.3	73.7±4.0	64.7±2.0	75.0±3.8	63.7±1.8	72.9±2.5	63.4±0.5	74.2±1.8
		obc_50	50.0±2.8	64.7±1.4	50.0±0.6	65.5±1.3	54.2±1.4	68.7±1.8	55.5±2.2	69.5±3.0
		obc_60	32.6±2.9	60.8±3.6	33.4±3.8	58.2±2.2	32.1±1.8	57.1±2.2	36.3±3.0	58.9±3.1

Table 30: TGAS with directions built from different numbers of samples per class. **Acc** is accuracy (%) and **Term** the share of generations that terminated (%). **n50**, **n75** and **n100** use 50, 75 and 100 samples per class; **Full** is the direction used in the main experiments. We use the same α values as in the main experiments for each configuration.

F.6 Generalization Across Pruning Methods

Table 32 shows that the termination failure is not specific to one pruning method. Under both RAC (SparseGPT) and C4-calibrated Wanda, the termination rate of Base drops monotonically as sparsity increases, for every model and benchmark. For DeepSeek-R1-Distill-Qwen-7B pruned with RAC, AIME-25 termination falls from 79.2% at 10% sparsity to 32.5% at 50% and 1.7% at 60%, and LiveCodeBench termination falls from 85.0% to 37.4% and 2.4%. The same trend holds for DeepSeek-R1-Distill-Qwen-14B with RAC and for Wanda-pruned checkpoints. Accuracy degrades alongside termination in all cases.

TGAS remains effective under both pruners: it raises termination over Base on 16 of 18 checkpoint-benchmark pairs and improves accuracy on all 18, achieving the best accuracy on 15. The gains are largest where termination has collapsed, e.g., on DeepSeek-R1-Distill-Qwen-7B with RAC at 60% sparsity, LiveCodeBench termination rises from 2.4% to 21.8% and accuracy from 2.4% to 13.4%, and on DeepSeek-R1-Distill-Qwen-14B with RAC at 60% sparsity, AIME-25 termination rises from 21.7% to 47.5% and LiveCodeBench accuracy from 6.1% to 18.2%. Under Wanda, TGAS improves DeepSeek-R1-Distill-Qwen-7B at 50% sparsity from 47.5% to 64.2% termination on AIME-25 and from 31.8% to 39.2% accuracy on LiveCodeBench. These results indicate that TGAS recovers part of the accuracy lost to pruning by restoring termination, independent of the pruning method used.

Model	SEAL	Layer mean	Unit-avg
DS-Qwen-1.5B	dense	0.172	0.108
	obc_50	0.245	0.269
DS-Qwen-7B	dense	0.283	0.187
	obc_50	0.306	0.265
DS-Qwen-14B	dense	0.269	0.221
	obc_50	0.244	0.214
DS-Llama-8B	dense	0.163	0.153
OT3-7B	dense	0.177	0.180
OT3-1.5B	dense	0.146	0.157

Table 31: Cosine between the steering directions of SEAL and TGAS, averaged over decoder layers. **Layer mean**: cosine of the unnormalised layer averages; **Unit-avg**: each layer normalised before averaging.

F.7 Cross-Testing with Different Directions

Sensitivity to the Sparsity Level of the Direction. TGAS constructs the direction at one sparsity level and reuses it across checkpoints. To test this choice, we construct additional directions at 40% and 60% sparsity for DS-Qwen-1.5B (default: 50%) and at 40% and 50% for DS-Qwen-7B (default: 60%), following Section 4.1, and evaluate them on the dense, 10%, 50% and 60% checkpoints, selecting α using the same procedure on a held-out OpenR1-Math split. As shown in Figure 11, all directions perform comparably: averaged over the eight benchmark-checkpoint settings, accuracy varies by at most 0.6 points across directions of the same model (20.4–20.9% for DS-Qwen-1.5B and 40.1–40.7% for DS-Qwen-7B), and termination by at most 4.1 points. Larger differences in individual settings appear only on AIME-25 (up to about 9 points for DS-Qwen-7B), whose 30 problems give standard deviations of 8–13 points. TGAS is therefore not sensitive to the sparsity level at which the direction is constructed, consistent with Section 4.2.

F.8 Comparison with SEAL

TGAS vs SEAL direction similarities. Table 31 compares TGAS and SEAL directions across layers and model families. All nine comparisons show weak positive alignment, with cosine similarities of 0.146–0.306 under Layer-mean averaging and 0.108–0.269 under Unit-avg averaging. Although alignment is slightly stronger for Qwen-7B and 14B, the results indicate that the two methods share only a limited directional component and are largely distinct.

SEAL with a Direction from the 50% Checkpoint. To test whether SEAL is limited by building its direction on the dense model, we also construct its vector on the 50% checkpoint and tune it in the same way as the dense vector, selecting the steering layer and strength on the held-out OpenR1-Math split (Table 34). Averaged over the 16 settings of DS-Qwen-1.5B and 7B on AIME-25 and LiveCodeBench, accuracy changes only from 26.9% to 27.6%, with similar termination (65.9% and 64.9%). Both remain below TGAS (30.2% accuracy at 65.8% termination), which achieves the best accuracy in 14 of the 15 settings where any method scores above zero. The SEAL direction is thus stable with respect to the checkpoint it is built on, but steering away from reflection recovers less accuracy than steering toward termination.

F.9 Comparison with Plan-CoT

Plan-CoT (Sheng et al., 2025) shows that LRMs plan their reasoning length before generation: question activations at the <think> position encode a direction whose magnitude scales with problem difficulty and the number of reasoning tokens. The direction is extracted by difference-in-means between hard and easy questions,

$$\mathbf{r}_\ell = \frac{1}{4}\sum_{i=2}^{5}(\boldsymbol{\mu}_{\ell,i} - \boldsymbol{\mu}_{\ell,1})$$

where $\boldsymbol{\mu}_{\ell,i}$ is the mean layer-ℓ activation at <think> over MATH questions of difficulty level i, and is applied as

$$\mathbf{h}_\ell \leftarrow \mathbf{h}_\ell + \lambda\,\mathbf{r}_\ell$$

at every layer and token position; $\lambda < 0$ shortens reasoning by raising the </think> logit. Plan-CoT therefore differs from TGAS in what the direction contrasts: it captures pre-generation difficulty planning, whereas TGAS contrasts post-branch continuations that terminate with those that loop.

Following the original setup, we construct the Plan-CoT vector on MATH-500 questions grouped by difficulty level, either from the dense model (Plan-CoT$_{\text{dense}}$) or from the 50%-sparse checkpoint (Plan-CoT$_{50\%}$), and apply it with negative coefficients to counter non-termination. We search $\lambda \in \{-0.5, -0.4, -0.3, -0.2, -0.1\}$ and report the best value per checkpoint and benchmark. The selected Plan-CoT coefficient is either -0.1 or -0.2, while the selected TGAS strength ranges from $+0.5$ to $+1.5$.

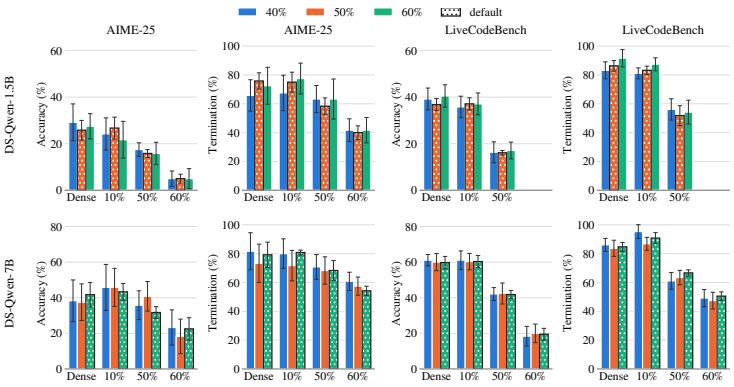

Figure 11: Accuracy and termination rate of DS-Qwen-1.5B and 7B when steered with directions constructed at different sparsity ratios. Error bars show the standard deviation over runs. The three directions give comparable results at every evaluated checkpoint.

		Acc (pooled)		Acc (best ckpt)		
Model	**Bench**	Δ	p	**Ckpt**	Δ	p
DS-Qwen-7B	AIME-25	+1.5	0.171	obc_10	+2.9	0.261
	LCB	**+1.5**	**0.016**	obc_60	**+4.4**	**0.003**
OT3-7B	AIME-25	**+2.1**	**0.046**	obc_50	**+5.4**	**0.002**
	LCB	**+1.1**	**0.026**	obc_60	**+2.3**	**0.030**

Table 33: Significance of TGAS accuracy gains over the base model (K=16). Δ: TGAS minus base (points); p: two-sided problem-level sign-flip test. Pooled: dense, 10%, 50% and 60% checkpoints combined. Best ckpt: largest gain, ties broken by lower p. Bold: $p < 0.05$.

Table 35 compares the TGAS and Plan-CoT on DeepSeek-R1-Distill-Qwen-7B and OpenThinker3-7B. In accuracy, TGAS is best or tied in 14 of 16 checkpoint benchmark pairs. TGAS also achieves the highest termination rate in 14 of 16 pairs over AIME-25 and LCB, with the largest gains at pruned checkpoints: on DeepSeek-R1-Distill-Qwen-7B, the termination rate of AIME-25 reaches 68.3% vs. 54.2% at 50% sparsity and 54.2% vs. 45.0% at 60% sparsity. The Plan-CoT vector built from the 50%-sparse checkpoint achieves comparable results to that from the dense model, but still underperforms TGAS.

F.10 Statistical Significance of TGAS Gains

Significance testing. We evaluate TGAS against the base model on AIME-25 and LiveCodeBench (LCB) using DS-Qwen-7B and OT3-7B, with $K = 16$ samples per problem. The steering strength α is fixed using a held-out OpenR1-Math split. We assess significance using a two-sided problem-level sign-flip permutation test with 20,000 permutations. Pooled results aggregate the dense, OBC at 10%, 50% and 60% sparsity checkpoints.

TGAS yields significant gains across model-benchmark pairs. Table 33 shows significant pooled accuracy gains in three of four model–benchmark pairs: DS-Qwen-7B on LCB (+1.5 points, $p = 0.016$) and OT3-7B on AIME-25 (+2.1, $p = 0.046$) and LCB (+1.1, $p = 0.026$). The largest checkpoint-level gains occur at pruned checkpoints, reaching +4.4 points for DS-Qwen-7B at 60% sparsity on LCB ($p = 0.003$) and +5.4 for OT3-7B at 50% sparsity on AIME-25 ($p = 0.002$). The pooled gain for DS-Qwen-7B on AIME-25 is not statistically significant (+1.5, $p = 0.171$). For this pair, TGAS primarily affects termination behavior rather than accuracy (Table 23).

TGAS reduces generation cost without sacrificing accuracy. Shorter generations are only useful if they do not come at the cost of accuracy. Across six models, four checkpoints and three benchmarks, TGAS matches or exceeds the base model in 64 of 72 settings and raises mean accuracy from 46.5% to 49.1% (Tables 23 and 24). Both effects grow with sparsity: from the dense model to 60% sparsity, the mean accuracy gain rises from +1.6 to +3.7 points, while the token reduction rises from 4.4% to 12.3%, or 1,423 tokens per generation (Table 8). As pruning makes non-termination more frequent, steering towards termination removes more wasted tokens and recovers more answers otherwise lost to the budget. Overall, TGAS makes pruned LRMs cheaper to run while keeping accuracy at or above the base model.

Model	Pruner	Ckpt	Method	AIME-25 Acc	AIME-25 Term	LiveCodeBench Acc	LiveCodeBench Term
Deepseek-R1-Distill -Qwen-7B	RAC	10%	Base	35.8±5.0	79.2±4.2	56.3±1.8	85.0±3.8
			No-Repeat	38.3±5.8	80.0±5.4	47.4±3.1	85.8±1.1
			Forced-Answer	33.3±4.7	—	53.7±1.2	—
			TGAS	**39.2±3.2**	85.0±6.9	**58.7±2.6**	89.7±3.0
		50%	Base	24.2±5.7	32.5±5.7	32.6±1.7	37.4±2.2
			No-Repeat	25.0±3.3	79.2±7.9	27.9±2.8	80.0±3.6
			Forced-Answer	23.3±3.8	—	**37.1±1.8**	—
			TGAS	**25.8±1.7**	59.2±7.4	34.5±4.0	47.9±1.1
		60%	Base	0.0±0.0	1.7±1.9	2.4±0.5	2.4±0.5
			No-Repeat	5.8±3.2	46.7±11.5	10.5±2.3	58.9±3.7
			Forced-Answer	**10.0±4.7**	—	6.1±2.3	—
			TGAS	6.7±5.4	10.8±7.9	**13.4±2.3**	21.8±2.3
Deepseek-R1-Distill -Qwen-14B	RAC	50%	Base	26.7±2.7	70.0±4.7	48.7±2.0	66.1±1.8
			No-Repeat	28.3±3.3	70.8±6.9	42.6±2.8	75.0±1.8
			Forced-Answer	24.2±4.2	—	47.6±1.0	—
			TGAS	**30.0±2.7**	68.3±4.3	**55.8±4.1**	68.7±3.3
		60%	Base	3.3±2.7	21.7±6.4	6.1±1.0	8.2±1.3
			No-Repeat	3.3±2.7	45.0±7.9	18.7±2.3	61.8±2.9
			Forced-Answer	3.3±2.7	—	**21.1±1.9**	—
			TGAS	**5.8±6.9**	47.5±8.8	18.2±2.3	26.8±1.4
Deepseek-R1-Distill -Qwen-7B	C4 Wanda	40%	Base	34.2±3.2	67.5±6.9	49.5±1.7	73.7±3.5
			No-Repeat	31.7±4.3	75.8±3.2	41.6±4.6	77.9±1.2
			Forced-Answer	30.0±4.7	—	49.7±2.0	—
			TGAS	**35.0±1.9**	75.8±5.0	**52.1±0.6**	81.6±4.7
		50%	Base	19.2±1.7	47.5±9.2	31.8±5.3	52.9±5.6
			No-Repeat	15.8±3.2	61.7±1.9	26.6±4.7	67.4±2.4
			Forced-Answer	19.2±1.7	—	34.7±4.3	—
			TGAS	**21.7±4.3**	64.2±7.4	**39.2±4.0**	58.7±3.6
		60%	Base	0.0±0.0	0.0±0.0	0.3±0.5	1.6±1.1
			No-Repeat	0.0±0.0	9.2±3.2	0.5±0.6	11.1±1.4
			Forced-Answer	0.0±0.0	—	0.3±0.5	—
			TGAS	**1.7±1.9**	1.7±3.3	**2.1±0.9**	14.5±1.6
OpenThinker3 -7B	C4 Wanda	40%	Base	37.5±5.7	53.3±8.2	50.5±1.9	64.2±2.9
			No-Repeat	39.2±6.9	58.3±6.9	45.5±4.1	65.3±2.1
			Forced-Answer	30.8±3.2	—	46.6±2.0	—
			TGAS	**40.0±2.7**	53.3±2.7	**55.0±2.0**	67.1±1.8
		50%	Base	17.5±3.2	40.8±5.7	29.7±2.0	58.9±0.9
			No-Repeat	18.3±4.3	44.2±5.7	24.2±4.8	54.7±4.5
			Forced-Answer	15.0±4.3	—	31.3±3.8	—
			TGAS	**20.0±3.8**	41.7±6.4	**31.6±3.6**	60.5±3.9

Table 32: Training-free remedies on checkpoints pruned with SparseGPT-pruned LRMs with RAC calibration and with C4 calibration (Wanda prune). **Acc** is accuracy (%) and **Term** the share of generations that terminated (%). Forced-Answer's **Term** is not reported, since forcing ends every generation by design. Best accuracy per checkpoint and dataset in bold.

		AIME-25						LiveCodeBench					
		SEAL (dense)		SEAL (50%)		TGAS		SEAL (dense)		SEAL (50%)		TGAS	
Model	Ckpt	Acc	Term	Acc	Term	Acc	Term	Acc	Term	Acc	Term	Acc	Term
DS-Qwen-1.5B	dense	18.3±1.9	76.7±6.1	19.5±2.9	75.8±8.3	**25.8±4.2**	75.8±5.7	35.0±3.7	84.7±3.3	35.5±2.2	86.6±3.0	**36.8±2.6**	86.3±3.7
	obc_10	21.7±4.3	80.8±7.4	25.0±6.9	76.7±8.6	**26.7±4.7**	75.0±6.9	34.5±4.6	85.5±3.0	33.4±3.6	81.3±2.2	**37.1±2.6**	83.2±2.9
	obc_50	15.8±3.2	60.0±6.7	**16.7±2.7**	57.5±5.7	15.8±1.7	58.3±5.8	13.2±3.5	46.3±4.0	15.3±1.8	45.8±5.0	**16.1±1.0**	51.8±6.9
	obc_60	**5.0±6.4**	27.5±7.4	3.3±2.7	24.2±7.4	**5.0±1.9**	40.0±4.7	0.0±0.0	0.0±0.0	0.0±0.0	0.0±0.0	0.0±0.0	2.4±1.8
DS-Qwen-7B	dense	36.7±2.7	80.8±3.2	38.3±1.9	84.2±6.9	**41.7±6.9**	79.2±8.8	57.9±1.2	90.0±2.6	56.3±2.5	90.5±1.5	**59.7±2.2**	89.2±3.0
	obc_10	32.5±10.7	80.8±5.7	32.5±5.7	84.2±1.7	**43.3±4.7**	80.8±1.7	55.8±2.6	86.1±1.8	56.3±0.6	86.8±2.2	**60.3±3.4**	90.8±3.8
	obc_50	29.2±5.7	70.8±5.0	30.0±2.7	70.0±4.7	**31.7±3.3**	68.3±6.9	40.3±1.6	75.0±3.9	40.4±1.8	70.8±1.6	**41.8±2.3**	66.6±2.2
	obc_60	16.7±4.7	60.8±7.9	21.7±4.3	56.7±0.0	**22.5±3.2**	54.2±3.2	18.4±1.4	48.4±2.7	16.8±3.4	46.8±8.2	**19.5±3.3**	50.5±3.0

Table 34: SEAL with steering vectors built on the dense model or on the 50% checkpoint, compared with TGAS. **Acc**: accuracy (%); **Term**: termination rate (%). Best accuracy per checkpoint and benchmark in bold.

Model	Sparsity	AIME-25			LiveCodeBench		
		TGAS	Plan-CoT$_{\text{dense}}$	Plan-CoT$_{50\%}$	TGAS	Plan-CoT$_{\text{dense}}$	Plan-CoT$_{50\%}$
Accuracy (%)							
DS-Qwen-7B	Dense	$\mathbf{41.7}_{\pm 6.9}$ (+1)	$37.5_{\pm 3.2}$ (-0.2)	$39.2_{\pm 5.0}$ (-0.2)	$59.7_{\pm 2.2}$ (+1)	$58.7_{\pm 2.8}$ (-0.1)	$\mathbf{60.3}_{\pm 3.0}$ (-0.2)
	10%	$\mathbf{43.3}_{\pm 4.7}$ (+1.25)	$40.0_{\pm 2.7}$ (-0.1)	$38.3_{\pm 4.3}$ (-0.2)	$\mathbf{60.3}_{\pm 3.4}$ (+1.25)	$\mathbf{60.3}_{\pm 2.3}$ (-0.1)	$57.9_{\pm 3.5}$ (-0.2)
	50%	$\mathbf{31.7}_{\pm 3.3}$ (+0.5)	$29.2_{\pm 3.2}$ (-0.1)	$\mathbf{31.7}_{\pm 7.9}$ (-0.2)	$\mathbf{41.8}_{\pm 2.3}$ (+0.5)	$40.3_{\pm 1.3}$ (-0.1)	$40.0_{\pm 3.1}$ (-0.2)
	60%	$22.5_{\pm 3.2}$ (+1.25)	$22.5_{\pm 3.2}$ (-0.1)	$\mathbf{25.0}_{\pm 1.9}$ (-0.1)	$\mathbf{19.5}_{\pm 3.3}$ (+1.25)	$15.5_{\pm 1.8}$ (-0.2)	$13.7_{\pm 1.2}$ (-0.2)
OT3-7B	Dense	$\mathbf{45.8}_{\pm 4.2}$ (+1.5)	$43.3_{\pm 6.1}$ (-0.2)	$45.0_{\pm 6.4}$ (-0.2)	$\mathbf{64.7}_{\pm 1.4}$ (+1.5)	$62.9_{\pm 2.5}$ (-0.1)	$62.1_{\pm 2.6}$ (-0.1)
	10%	$\mathbf{48.3}_{\pm 5.8}$ (+1.5)	$44.2_{\pm 3.2}$ (-0.1)	$45.8_{\pm 4.2}$ (-0.1)	$\mathbf{63.4}_{\pm 0.5}$ (+1.5)	$62.1_{\pm 1.5}$ (-0.1)	$62.1_{\pm 2.3}$ (-0.1)
	50%	$\mathbf{43.3}_{\pm 8.2}$ (+1.25)	$39.2_{\pm 7.4}$ (-0.2)	$39.2_{\pm 1.7}$ (-0.1)	$\mathbf{55.5}_{\pm 2.2}$ (+1.25)	$52.1_{\pm 2.0}$ (-0.1)	$52.9_{\pm 2.8}$ (-0.1)
	60%	$\mathbf{29.2}_{\pm 1.7}$ (+1.5)	$26.7_{\pm 2.7}$ (-0.2)	$24.2_{\pm 5.7}$ (-0.2)	$\mathbf{36.3}_{\pm 3.0}$ (+1.5)	$34.5_{\pm 3.5}$ (-0.1)	$34.5_{\pm 2.8}$ (-0.1)
Termination rate (%)							
DS-Qwen-7B	Dense	$\mathbf{79.2}_{\pm 8.8}$ (+1)	$75.0_{\pm 5.8}$ (-0.2)	$77.5_{\pm 1.7}$ (-0.2)	$\mathbf{89.2}_{\pm 3.0}$ (+1)	$86.6_{\pm 2.9}$ (-0.1)	$82.6_{\pm 2.5}$ (-0.2)
	10%	$\mathbf{80.8}_{\pm 1.7}$ (+1.25)	$77.5_{\pm 5.7}$ (-0.1)	$75.8_{\pm 6.9}$ (-0.2)	$\mathbf{90.8}_{\pm 3.8}$ (+1.25)	$85.0_{\pm 0.5}$ (-0.1)	$81.8_{\pm 4.9}$ (-0.2)
	50%	$\mathbf{68.3}_{\pm 6.9}$ (+0.5)	$54.2_{\pm 4.2}$ (-0.1)	$52.5_{\pm 6.9}$ (-0.2)	$66.6_{\pm 2.2}$ (+0.5)	$\mathbf{67.6}_{\pm 2.0}$ (-0.1)	$63.4_{\pm 3.4}$ (-0.2)
	60%	$\mathbf{54.2}_{\pm 3.2}$ (+1.25)	$45.0_{\pm 8.8}$ (-0.1)	$44.2_{\pm 5.7}$ (-0.1)	$\mathbf{50.5}_{\pm 3.0}$ (+1.25)	$47.4_{\pm 3.1}$ (-0.2)	$42.6_{\pm 3.0}$ (-0.2)
OT3-7B	Dense	$\mathbf{54.2}_{\pm 3.2}$ (+1.5)	$47.5_{\pm 3.2}$ (-0.2)	$50.0_{\pm 7.2}$ (-0.2)	$\mathbf{73.9}_{\pm 1.8}$ (+1.5)	$72.4_{\pm 3.0}$ (-0.1)	$70.3_{\pm 2.9}$ (-0.1)
	10%	$\mathbf{56.7}_{\pm 5.4}$ (+1.5)	$51.7_{\pm 6.4}$ (-0.1)	$50.8_{\pm 7.4}$ (-0.1)	$\mathbf{74.2}_{\pm 1.8}$ (+1.5)	$71.8_{\pm 1.0}$ (-0.1)	$71.6_{\pm 3.6}$ (-0.1)
	50%	$\mathbf{52.5}_{\pm 9.6}$ (+1.25)	$47.5_{\pm 8.8}$ (-0.2)	$49.2_{\pm 5.7}$ (-0.1)	$\mathbf{69.5}_{\pm 3.0}$ (+1.25)	$65.5_{\pm 2.6}$ (-0.1)	$66.1_{\pm 3.6}$ (-0.1)
	60%	$50.8_{\pm 6.9}$ (+1.5)	$\mathbf{51.7}_{\pm 6.4}$ (-0.2)	$48.3_{\pm 7.9}$ (-0.2)	$\mathbf{58.9}_{\pm 3.1}$ (+1.5)	$56.3_{\pm 3.0}$ (-0.1)	$56.8_{\pm 3.9}$ (-0.1)

Table 35: TGAS at its best steering strength α (shown in parentheses) compared with Plan-CoT at best steering coefficient, using a steering vector extracted from the dense model (Plan-CoT$_{\text{dense}}$) or from the 50%-sparse model (Plan-CoT$_{50\%}$). Best value per row and benchmark in bold.